



Computational inference time — all architectures
GPU: NVIDIA Tesla T4 (Google Colab) | N=10 runs per configuration | 3 warmup passes excluded

Total inference time per batch



Per-image inference time



■ Vision Transformer ■ Xception ■ DenseNet121 ■ ResNet50