


Graphical Abstract

Regression-Based Explainable Deep Learning for Estimating Hamiltonian Parameters from Magnetic Nanodot Images

Juan Sebastián Méndez-Rondón, María Isabel García-Quimbayo, Andrés Marino Álvarez-Meza, Jorge Iván Montes-Monsalve, Jose Darío Agudelo-Giraldo



Highlights

Regression-Based Explainable Deep Learning for Estimating Hamiltonian Parameters from Magnetic Nanodot Images

Juan Sebastián Méndez-Rondón, María Isabel García-Quimbayo, Andrés Marino Álvarez-Meza, Jorge Iván Montes-Monsalve, Jose Darío Agudelo-Giraldo

- Atomistic Hamiltonian generates relaxed 2D magnetic domain images for training.
- Deep learning regression recovers Hamiltonian parameters from nanodot images.
- Regression Activation Maps explain continuous prediction errors spatially.
- Our framework isolates degenerate textures linking black-box AI to magnetic physics.

Regression-Based Explainable Deep Learning for Estimating Hamiltonian Parameters from Magnetic Nanodot Images

Juan Sebastián Méndez-Rondón¹, María Isabel García-Quimbayo¹, Andrés Marino Álvarez-Meza¹, Jorge Iván Montes-Monsalve¹, Jose Darío Agudelo-Giraldo¹

^aAI-Lab, Signal Processing and Recognition Group, Universidad Nacional de Colombia, Sede Manizales, 170003 Manizales, Colombia

^bGrupo de Investigación en Física y Matemáticas con Énfasis en la Formación de Ingenieros, Departamento de Física y Matemáticas, Facultad de Ingeniería, Universidad Autónoma de Manizales, Antigua Estación del Ferrocarril, 170002 Manizales, Colombia

Abstract

Magnetic nanodots host topological spin textures whose stability is dictated by an extended Heisenberg Hamiltonian combining symmetric exchange, the Dzyaloshinskii–Moriya interaction, magnetocrystalline anisotropy, and the Zeeman field. Inferring these atomistic parameters from two-dimensional magnetization images underpins the design of functional magnetic materials, yet it remains an ill-posed inverse problem: under certain external conditions and transient states, configurations relax into noisy, weakly discriminative textures that render the dataset ill-conditioned. Deep-learning estimators act as opaque black boxes and rely on slow atomistic simulations. **We propose an explainable regression methodology whose novelty lies in data generation and interpretability: a matrix decomposition of the extended Hamiltonian enables GPU-accelerated Metropolis–Hastings Monte Carlo with simulated annealing, accelerating computation without altering the underlying physics, while Regression Activation Maps (RAMs), a novel error-based attribution metric, turn the network into a traceable gray-box model.** Trained on several configurations, convolutional and Vision Transformer regressors reliably recover the physical parameters from simulated nanodot images, while the higher-order exchange constants remain bounded by the intrinsic identifiability limits of the s_z projection. Comparing both families through their attention maps and accuracy-versus-size trade-off, RAMs isolate the chiral domain walls and skyrmion cores driving predictions.

Keywords: Explainable AI, Magnetic Nanodots, Hamiltonian Parameters, Regression Activation Maps, Deep Learning

1. Introduction

Understanding and controlling magnetic behavior at the nanometric scale is a central challenge in condensed matter physics and materials engineering [?], particularly relevant to spintronic devices, ultrahigh-density memories, and neuromorphic platforms based on topological magnetic textures [?]. At such length scales, surface effects and geometric confinement substantially modify magnetic properties relative to the bulk [?], enhancing surface anisotropy and thermal fluctuations that affect the stability of spin configurations [?]. These effects give rise to emergent states: ferromagnetic, helical, labyrinthine, and skyrmionic textures [?]; governed by competing microscopic interactions [?]. Magnetic nanodots are an ideal model platform, as their lateral confinement favors single-domain states, vortices, and skyrmion-like configurations controlled by the balance among symmetric exchange, anisotropy, the DMI, and external fields [?]. Inferring these parameters directly from magnetic domain images has thus become a critical inverse problem for the rational engineering of functional magnetic materials [?].

Atomistic Hamiltonian-based modeling provides the most rigorous description of the discrete spin interactions and boundary effects that dominate nanoscale magnetic systems. Unlike

continuum micromagnetic models, which can mask surface and finite-size contributions, the extended Heisenberg Hamiltonian explicitly captures the interactions that stabilize complex topological textures [?]. It combines the symmetric exchange ($\tilde{J}$ [meV]), promoting collinear alignment; the Dzyaloshinskii–Moriya interaction, DMI ($\tilde{K}_{DM}$ [meV]), favoring chiral textures; magnetocrystalline anisotropy ($\tilde{K}$ [meV]), selecting preferred orientations; and the Zeeman coupling to an external field ($\tilde{H}_{ex}$ [meV]) [?]. Their competition generates a highly multimodal energy landscape in which distinct parameter sets can yield nearly indistinguishable out-of-plane textures [?]. Compounded by the noise of external conditions and transient states during thermal relaxation, this degeneracy renders the inverse estimation intrinsically ill-conditioned, imposing a physical limit on identifiability that must be explicitly characterized [?].

The ambiguity of the inverse problem is further exacerbated by the limitations of current experimental and analytical methodologies. Advanced imaging techniques, including magnetic force microscopy, Lorentz transmission electron microscopy, and spin-polarized scanning tunneling microscopy [?], enable high-resolution visualization of magnetic textures, but they typically provide indirect or transformed measurements that are subject to modality-dependent distortions, noise, and finite spatial resolution [?]. Parameter estimation therefore relies heavily on empirical fitting or on comparison with synthetic magnetic configurations generated by atomistic spin simulations or con-

*amalvarezme@unal.edu.co

tinuum micromagnetic models [?]. Yet similarity measures defined in image space, such as mean squared error or structural similarity, are often inadequate for resolving topologically distinct states, since they primarily quantify visual resemblance rather than physical equivalence. Consequently, the optimization process remains structurally ambiguous, and different Hamiltonian parameter sets may produce magnetic domain patterns that are visually indistinguishable [?].

To circumvent the prohibitive cost of iterative simulation-based optimization, recent efforts have reformulated this inverse problem as a data-driven learning task using deep neural architectures [?]. Although supervised regression models allow fast parameter inference, they generally impose an implicit one-to-one mapping and therefore tend to collapse the intrinsic multimodality of degenerate magnetic systems into over-smoothed point predictions [?]. Methods designed to explicitly capture this uncertainty, such as Bayesian inference, diffusion models, cycle-consistent networks, and physics-informed neural networks, offer improved representational flexibility, but often introduce substantial computational overhead and pronounced optimization stiffness [?]. Beyond these computational challenges, a central limitation remains their opacity. In strongly degenerate parameter spaces, conventional black-box models offer little or no mechanistic understanding of how localized spatial features contribute to the inferred Hamiltonian parameters [?]. Moreover, although gradient-based attribution tools such as class activation maps are widely used in classification, analogous interpretability methods for continuous regression in highly multimodal physical systems remain underdeveloped [?]. This interpretability gap restricts the physical traceability required for robust and scientifically meaningful machine-learning-based inference.

Recent advances have demonstrated the potential of deep learning in resolving the inverse problem of magnetic Hamiltonian parameter estimation from domain images. For example, Kwon *et al.* [?] employed ResNet architectures to estimate DMI, anisotropy, and dipolar interactions from Monte Carlo-simulated spin configurations, while Kong *et al.* [?] utilized customized convolutional neural networks to extract exchange stiffness, effective anisotropy, and interfacial DMI from micromagnetic simulations. More recently, Lee *et al.* [?] demonstrated the applicability of fully connected networks for estimating interlayer and intralayer interactions in twisted van der Waals magnets. Despite their predictive accuracy, these pioneering approaches predominantly treat the neural network as a black box. Specifically, while some works analyze intermediate feature maps [?] or perform statistical error analysis [?], they lack spatially resolved attribution to explicitly identify which morphological features of the spin textures drive specific parameter predictions. Also, physical identifiability limits, such as degenerate spin configurations arising from competing symmetric interactions, are often encountered as empirical estimation errors rather than explicitly quantified boundaries.

We propose an explainable deep learning framework grounded in an atomistic extended Heisenberg Hamiltonian, coupling Metropolis–Hastings Monte Carlo with simulated annealing to generate thermodynamically relaxed magnetization fields that

are then fed to a regression-based architecture for inverse parameter estimation. Our main contributions are threefold:

- We design a tensor decomposition of the extended Hamiltonian that recasts the atomistic energy as linear-algebra operations, enabling a GPU-accelerated Monte Carlo simulation, without altering the underlying physics.
- We propose the Regression Activation Maps (RAMs), a novel error-based attribution method extending gradient activation mapping to continuous physical targets, which turns the regressor into a traceable gray-box model linking predictions to physically meaningful spin textures.
- We compare convolutional and Vision Transformer regressors, contrasting their attention and RAM maps as well as their accuracy relative to model size, and explicitly map the identifiability boundaries of the Hamiltonian parameters from the nanodot image, distinguishing recoverable parameters from degenerate ones.

To assess the effectiveness of the proposed framework, extensive experiments were performed using a custom-generated dataset of simulated magnetic domain images covering a broad range of Hamiltonian parameter configurations. The obtained results indicate that the proposed method enables accurate inverse prediction of a subset of the Hamiltonian parameters, most notably the DMI and symmetric exchange, while explicitly exposing the identifiability limits of the remaining terms under the s_z projection, with competitive performance across the evaluated convolutional and transformer models. Moreover, the qualitative and quantitative analyses based on RAMs indicate that the learned representations are associated with physically relevant spin topologies rather than dataset-specific noise or superficial correlations, thereby strengthening the interpretability and reliability of the proposed approach.

The remainder of the paper is organized as follows: Section ?? reviews related work; Section ??, the materials and methods; Section ??, the experimental setup; Section ??, the results; and Section ?? concludes with future directions.

2. Related Work

The inverse estimation of Hamiltonian parameters from magnetic domain configurations is widely recognized as a highly challenging problem due to the complex, nonlinear relationship between physical interactions and observable spin textures. From an experimental standpoint, researchers heavily rely on advanced indirect imaging techniques such as magnetic force microscopy, Lorentz transmission electron microscopy, and spin-polarized scanning tunneling microscopy [?]. While these modalities enable high-resolution visualization of magnetic domains, they inherently capture transformed or partial projections of the underlying three-dimensional spin configurations. Because each technique is governed by distinct physical interaction principles, experimental observations are frequently contaminated by modality-specific distortions, finite spatial resolution

limits, and heterogeneous noise, making direct inversion from these measurements generally infeasible [? ?].

Consequently, practical parameter estimation heavily depends on empirical fitting or direct comparison with synthetic data generated through numerical simulations [?]. Computationally the problem is most commonly addressed using physics-based forward simulations, specifically atomistic spin models governed by the extended Heisenberg Hamiltonian and continuum micromagnetic models [? ? ?]. However, attempting to optimize the inverse mapping by minimizing pixel-level discrepancy functions, such as Mean Squared Error or Mean Absolute Error, frequently fails to capture physically meaningful topological differences or vital spin interactions [?]. Even structural metrics like the Structural Similarity Index fall short in this regard [?]. Ultimately, because the comparison between simulated and experimental data occurs strictly in the image space rather than the parameter space, the optimization landscape becomes highly ambiguous, allowing disparate parameter configurations to yield visually identical domain patterns [?].

To bypass the computationally prohibitive nature of iterative simulation-based optimization, recent advances have increasingly framed the inverse problem as a data-driven mapping utilizing deep learning architectures [?]. The most prevalent approach employs supervised regression models, including convolutional neural networks and transformer-based architectures, to approximate this mapping deterministically. While these models offer rapid inference and scalability, they inherently assume one-to-one mapping between images and parameters. In highly degenerate magnetic systems, this assumption causes the models to collapse the intrinsic multimodality of the parameter space into a single, often smoothed, point estimate, rendering them highly sensitive to dataset bias and sim-to-real mismatch [?].

To explicitly model uncertainty and capture this multimodality, probabilistic frameworks such as Bayesian inference methods, conditional variational autoencoders, and diffusion models have been introduced [?]. Despite their theoretical flexibility in sampling multiple plausible configurations, these generative and Bayesian approaches suffer from immense computational overhead, unstable training dynamics, and a distinct lack of explicit physical guarantees. In parallel, cycle-consistent frameworks and Physics-Informed Neural Networks attempt to constrain the solution space by enforcing coherence between forward and inverse mappings, or by directly embedding thermodynamic penalties into the loss function [? ?]. While physically plausible, these hybrid models introduce severe optimization stiffness and rely fundamentally on the accuracy of the underlying forward surrogate, which can dangerously propagate errors within strongly degenerate spin regimes.

Despite their growing predictive power, the fundamental opacity of deep neural networks remains a barrier to their adoption in condensed matter physics. In nanoscale magnetism, the competition among exchange, DMI, and anisotropy creates a densely degenerate parameter space where distinct configurations present nearly identical free energies [?], yet black-box architectures provide no mechanistic insight into how spatial features correlate with the predicted Hamiltonian constants [?], preventing researchers from telling whether a model has captured the under-

lying physics or merely memorized dataset artifacts. Moreover, while gradient-based attribution methods such as Class Activation Maps are widely used for discrete classification, their extension to continuous regression in multimodal physical systems remains largely underexplored [?]. This gap restricts the scientific traceability of such models, motivating architectures that couple high predictive accuracy with transparent, physics-informed spatial interpretations.

To overcome these limitations, our methodology departs from prior work along three concrete axes. First, whereas most existing approaches operate at the continuum micromagnetic scale, we adopt an atomistic description at the nanometric scale, grounded in a rigorous extended Heisenberg Hamiltonian that explicitly resolves the discrete spin interactions and the surface and finite-size effects governing magnetic nanodots. Second, in contrast to the predominantly black-box models used in this field, which expose no mechanism to trace how spatial features drive the inferred parameters, we introduce RAMs, an error-based attribution method that correlates localized spatial activations with the absolute prediction error of the continuous physical targets, turning the regressor into a transparent, physically traceable gray-box model. Third, rather than relying on conventional classical simulation software, we build the data-generation process from a matrix/tensor (algebraic) decomposition of the extended Hamiltonian that, coupled with GPU computing, substantially accelerates the Metropolis–Hastings Monte Carlo with simulated annealing without altering the underlying physics. Together, these choices couple physics-grounded, high-throughput simulation with an inherently explainable regression model, enabling the precise isolation of the degenerate magnetic textures, such as chiral domain walls and skyrmions, that drive specific parameter estimations, and establishing a reliable, transparent, and traceable pathway for analyzing complex nanoscale magnetic systems. These distinctions differentiate our framework from representative deep-learning approaches for magnetic Hamiltonian parameter estimation, e.g., [? ? ?], across several technical dimensions: the Hamiltonian terms modeled, the simulation method, the deep-learning architecture, the number of estimated parameters, the evaluation metrics, and, crucially, the interpretability mechanism.

3. Materials and Methods

3.1. Atomistic Hamiltonian Framework

The theoretical foundation for modeling interacting magnetic dipole moments originates from the seminal formulations of Ernst Ising and Werner Heisenberg [?]. Ising introduced a one-dimensional statistical mechanics model wherein localized spins were heavily constrained to discrete, scalar states (i.e., rigidly pointing “up” or “down”). While the Ising model provided early fundamental insights into macroscopic ferromagnetism and phase transitions, its strict uni-axial constraint renders it structurally insufficient for capturing the continuous rotational degrees of freedom required to model complex, multi-dimensional magnetic textures. This limitation was subsequently resolved by Heisenberg, who formulated a generalized model allowing

spin vectors to vary continuously in three-dimensional space with rotational symmetry. While the foundational Heisenberg model successfully captures isotropic exchange driven by the Pauli exclusion principle, modern nanoscale applications, such as the study of magnetic vortices and topologically protected skyrmions in confined nanodot geometries, necessitate an extended approach [?]. This extended framework rigorously incorporates symmetry-breaking relativistic spin-orbit coupling effects, primarily the DMI and magnetocrystalline anisotropy alongside standard Zeeman coupling [?].

Formally, let $\mathcal{V} = \{1, \dots, N\}$ denote the set of indices for the discrete magnetic lattice sites. The thermodynamic state of the system is characterized by the spatial distribution of localized magnetic moments, represented by the normalized spin vectors $\mathbf{s}_i \in \mathbb{R}^3$, s.t. $\|\mathbf{s}_i\|_2 = 1$, for $i \in \mathcal{V}$. In a generalized theoretical context, all pairwise symmetric and antisymmetric interactions as well as site-specific field and anisotropy alignments, can be collapsed into a unified geometric approach [?]. The total effective energy of the system $\mathcal{H} \in \mathbb{R}$ [meV] is thus defined by the following generalized quadratic form:

$$\mathcal{H}(\mathbf{S}|\boldsymbol{\theta}) = - \sum_{i=1}^N \sum_{j \in \mathcal{N}(i)} \langle \mathbf{s}_i, \tilde{\mathbf{M}}_{ij} \mathbf{s}_j \rangle - \sum_{i=1}^N \langle \tilde{\mathbf{h}}_i, \mathbf{s}_i \rangle - \sum_{i=1}^N \langle \mathbf{s}_i, \tilde{\mathbf{K}}_i \mathbf{s}_i \rangle, \quad (1)$$

where $\mathcal{N}(i)$ defines the set of interacting nearest neighbors for a given site i , and the matrix $\mathbf{S} \in \mathbb{R}^{N \times 3}$ gathers the continuous spin vectors across the lattice with $\mathbf{s}_i \in \mathbf{S}$. Here, $\boldsymbol{\theta} = \{\tilde{\mathbf{M}}_{ij}, \tilde{\mathbf{h}}_i, \tilde{\mathbf{K}}_i\}$, where $\tilde{\mathbf{M}}_{ij} \in \mathbb{R}^{3 \times 3}$ represents the complete bilinear interaction matrix coupling neighboring spins, $\tilde{\mathbf{h}}_i \in \mathbb{R}^3$ accounts for linear external field interactions, and $\tilde{\mathbf{K}}_i \in \mathbb{R}^{3 \times 3}$ is the anisotropy tensor governing site-specific directional preferences. Casting every interaction as dense linear-algebra operations acting on the spin matrix $\mathbf{S}$ is a deliberate modeling choice: it recasts the atomistic energy as batched matrix and tensor products rather than the sequential site-by-site loops of conventional atomistic codes, which is precisely what enables the GPU-accelerated evaluation and sampling detailed in Section ??, without altering the underlying physics.

Afterward, to recover the isolated physical phenomena and define an extended target scalar parameter space $\boldsymbol{\theta} = [\tilde{J}_1, \tilde{J}_2, \tilde{J}_3, \tilde{J}_4, \tilde{K}_{DM}, \tilde{H}_{ex}, \tilde{K}_{an1}, \tilde{K}_{anS}]$, the generalized matrices must be mathematically decomposed. By factorizing these parameters, we can map each matrix operation to its specific energetic contribution:

Bilinear Coupling ($\tilde{\mathbf{M}}_{ij}$): A real square matrix can be separated into its symmetric and skew-symmetric components. $\tilde{\mathbf{M}}_{ij}$ is parameterized by factorizing the isotropic exchange $\tilde{J}_r$ for the r -th neighbor shell and the DMI magnitude $\tilde{K}_{DM}$ along a predefined spatial direction vector $\hat{\mathbf{d}}_{ij} \in \mathbb{R}^3$ ($\|\hat{\mathbf{d}}_{ij}\|_2 = 1$), yielding:

$$\tilde{\mathbf{M}}_{ij}(\tilde{J}_r, \tilde{K}_{DM}) = \tilde{J}_r \mathbf{I} - \tilde{K}_{DM} [\hat{\mathbf{d}}_{ij}]_{\times}, \quad (2)$$

where $\mathbf{I} \in \mathbb{R}^{3 \times 3}$ represents the identity matrix and $[\hat{\mathbf{d}}_{ij}]_{\times} \in \mathbb{R}^{3 \times 3}$ denotes the skew-symmetric matrix representation of the cross-product. The symmetric exchange is modeled across four neighbor shells $r \in \{1, 2, 3, 4\}$, where $\mathcal{N}_r(i)$ denotes the set of neighbors at shell r around site i , ordered by increasing interatomic

distance ($d_1 < d_2 < d_3 < d_4$). To resolve the fundamental scale ambiguity of the inverse problem (whereby multiplying all parameters by a common factor leaves the observable spin configuration unchanged) the first-shell exchange is fixed as the energy reference unit, $\tilde{J}_1 \in \mathbb{R}$ in meV, while $\tilde{J}_2, \tilde{J}_3, \tilde{J}_4 \in \mathbb{R}$ are free parameters (also in meV). The DMI acts exclusively between first-shell neighbors and is indexed over $\mathcal{N}_1(i)$.

Zeeman Interaction ($\tilde{\mathbf{h}}_i$): The linear field term is factorized by extracting the scalar field magnitude $\tilde{H}_{ex}$ and a fixed normalized direction vector $\hat{\mathbf{h}}$, with $\|\hat{\mathbf{h}}\|_2 = 1$, of the externally applied magnetic field, yielding $\tilde{\mathbf{h}}_i = \tilde{H}_{ex} \hat{\mathbf{h}} \in \mathbb{R}^3$. The orientation $\hat{\mathbf{h}}$ is held constant within each simulation but is selected among the principal crystallographic axes, either in-plane ($\hat{\mathbf{e}}_x$) or out-of-plane ($\hat{\mathbf{e}}_z$), to stabilize distinct families of topological textures. Fixing this direction eliminates the orientational degree of freedom of $\tilde{\mathbf{h}}_i$, ensuring that $\tilde{H}_{ex}$ alone serves as the free parameter governing the Zeeman contribution.

Magnetocrystalline Anisotropy ($\tilde{\mathbf{K}}_i$): Modeling cubic magnetocrystalline anisotropy requires factorizing a site-dependent scalar constant $\tilde{K}_i^{eff}$ to scale a quartic interaction. The principal crystallographic axes of the nanodot are defined by the orthonormal basis vectors $\mathbf{e}_{\hat{x}}, \mathbf{e}_{\hat{y}}, \mathbf{e}_{\hat{z}} \in \mathbf{E}$, with $\mathbf{E} \in \mathbb{R}^{3 \times 3}$, and the directional cosines are given by $s_{im} = \langle \mathbf{e}_m, \mathbf{s}_i \rangle$ ($\forall m \in \{\hat{x}, \hat{y}, \hat{z}\}$). Because atoms at the nanodot surface experience broken inversion symmetry (having fewer neighbors on the exterior side), their anisotropy differs fundamentally from that of interior atoms. The lattice is therefore partitioned into two disjoint subsets: the bulk region $\mathcal{V}_{bulk} \subset \mathcal{V}$, comprising interior sites with complete coordination shells, and the surface region $\mathcal{V}_{surf} \subset \mathcal{V}$, comprising boundary sites exposed to the geometric confinement, with $\mathcal{V}_{bulk} \cup \mathcal{V}_{surf} = \mathcal{V}$ and $\mathcal{V}_{bulk} \cap \mathcal{V}_{surf} = \emptyset$. This yields the site-dependent effective anisotropy constant:

$$\tilde{K}_i^{eff} = \begin{cases} \tilde{K}_{an1} & \text{if } i \in \mathcal{V}_{bulk}, \\ \tilde{K}_{anS} & \text{if } i \in \mathcal{V}_{surf}, \end{cases} \quad (3)$$

where $\tilde{K}_{an1}, \tilde{K}_{anS} \in \mathbb{R}$ are scalar constants in meV per atom [?].

Substituting these factorized relationships back into Equation ?? yields the following explicit formulation for the total extended Hamiltonian:

$$\begin{aligned} \mathcal{H}(\mathbf{S}|\hat{\mathbf{d}}_{ij}, \hat{\mathbf{h}}; \boldsymbol{\theta}) = & - \sum_{r=1}^4 \tilde{J}_r \sum_{i=1}^N \sum_{j \in \mathcal{N}_r(i)} \langle \mathbf{s}_i, \mathbf{s}_j \rangle \\ & + \tilde{K}_{DM} \sum_{i=1}^N \sum_{j \in \mathcal{N}_1(i)} \langle \hat{\mathbf{d}}_{ij}, (\mathbf{s}_i \times \mathbf{s}_j) \rangle \\ & - \tilde{H}_{ex} \sum_{i=1}^N \langle \hat{\mathbf{h}}, \mathbf{s}_i \rangle \\ & + \frac{1}{s_o^4} \sum_{i=1}^N \tilde{K}_i^{eff} (s_{ix}^2 s_{iy}^2 + s_{ix}^2 s_{iz}^2 + s_{iy}^2 s_{iz}^2), \end{aligned} \quad (4)$$

with $\tilde{J}_1 = 1.0$ meV fixed as the energy reference unit. The cubic anisotropy form in the last term follows directly from expanding $\frac{1}{2}(1 - \|\mathbf{s}_i\|_4^4)$ under the L_2 normalization constraint $\|\mathbf{s}_i\|_2 = 1$,

which yields $s_{ix}^2 s_{iy}^2 + s_{iy}^2 s_{iz}^2 + s_{iz}^2 s_{ix}^2$, normalized by s_o^4 to maintain416 dimensional consistency across spin magnitudes.417

Physically, the extended parameter space θ actively governs418 the competing energetic contributions that shape the magnetic419 energy landscape. Each $\tilde{J}_r$ dictates the collinear alignment ten420 dency between spin pairs separated by the corresponding dis421 tance $[?]$: positive values ($\tilde{J}_r > 0$) favor parallel, ferromagnetic422 alignment, while negative values ($\tilde{J}_r < 0$) drive antiparallel, anti423 ferromagnetic order. This multi-shell formulation is essential in424 confined nanodot geometries, where longer-range exchange path425 ways, mediated through intermediate atomic orbitals, contribute426 non-negligibly to the global energetic competition and can se427 lectively stabilize or destabilize complex topological textures428 beyond what first-shell interactions alone would predict $[?]$ 429 Among these constants, $\tilde{J}_1$ plays a privileged role: it corresponds430 to the strongest and most spatially direct exchange pathway, cou431 pling each spin to its nearest neighbors in the cubic lattice and432 establishing the dominant energy scale of the magnetic system433 For this reason, $\tilde{J}_1$ is fixed to 1.0 meV throughout all simulations434 serving as the absolute energy reference unit against which all435 remaining interaction strengths are expressed. The remaining436 constants $\tilde{J}_2$, $\tilde{J}_3$, and $\tilde{J}_4$ are free parameters whose magnitudes437 decay with increasing shell distance, consistent with the expo438 nential attenuation of quantum mechanical exchange integrals439 with interatomic separation $[?]$.440

The DMI magnitude, represented by $\tilde{K}_{DM}$, is responsible for441 driving chiral, non-collinear spin winding in systems exhibit442 ing broken spatial inversion symmetry. By actively competing443 with the dominant first-shell exchange $\tilde{J}_1$, this antisymmetric444 contribution stabilizes complex topological textures, such as445 chiral domain walls and skyrmions. It acts exclusively between446 first-shell neighbors, as longer-range antisymmetric exchange447 pathways are negligibly small compared to the dominant spin448 orbit coupling at nearest-neighbor interfaces. Measured in meV449 per interacting spin pair, its magnitude is fundamentally depen450 dent on the strength of the spin-orbit coupling at the material's451 interfaces and is naturally expressed as a fraction of the first452 shell exchange energy $\tilde{J}_1$ $[?]$; the specific sampled range and its453 physical justification are detailed in Section ??454

The Zeeman coupling energy, parameterized as $\tilde{H}_{ex}$, quanti455 fies the interaction between the localized atomic spins and the456 externally applied magnetic field $[?]$. It systematically penalizes457 magnetic misalignment with the fixed field axis $\hat{\mathbf{h}}$, frequently458 serving as the primary thermodynamic driving force for macro459 scopic magnetic phase transitions. This parameter is expressed460 in meV per atom, representing the effective energy equivalence461 ($g\mu_B B$) that transitions a macroscopic magnetic field B measured462 in Tesla into the atomistic energy regime; with $g \approx 2$ this cor463 responds to ≈ 0.116 meV per Tesla, so that the sampled range464 maps directly onto laboratory-accessible fields. Also, fixing the465 field direction along $\hat{\mathbf{h}}$ reduces the dimensionality of the inverse466 problem by eliminating the orientational degree of freedom of467 $\hat{\mathbf{h}}$, ensuring that $\tilde{H}_{ex}$ alone governs the Zeeman contribution as468 a single identifiable scalar. Notably, when the field is applied469 out-of-plane ($\hat{\mathbf{h}} = \hat{\mathbf{e}}_z$), the Zeeman term couples directly to the470 observable s_z component, which, as discussed below, renders471 $\tilde{H}_{ex}$ recoverable from the two-dimensional projection.

In turn, the magnetocrystalline anisotropy encodes the in4 intrinsic directional preference for magnetization relative to the underlying crystal lattice $[?]$. Arising from relativistic spin-orbit coupling, this term rigorously establishes the energetically favorable (easy) and unfavorable (hard) crystallographic axes for spin orientation within the nanodot. Both constants are measured in meV per atom and represent a relatively subtle symmetry-breaking energy compared to the dominant isotropic exchange. The bulk anisotropy constant $\tilde{K}_{an1}$ governs interior lattice sites $i \in \mathcal{V}_{bulk}$, where each atom is fully coordinated by neighbors in all spatial directions, yielding a locally symmetric energetic environment. Its magnitude ranges from near-zero in soft magnets to several meV per atom in strongly anisotropic interfacial systems $[?]$, with the sampled range specified in Section ?? The surface anisotropy constant $\tilde{K}_{anS}$, by contrast, governs boundary sites $i \in \mathcal{V}_{surf}$, where broken spatial inversion symmetry, arising from the absence of neighbors on the exterior side of the nanodot, produces an asymmetric local environment that intensifies the directional preference for spin orientation. Consequently, $|\tilde{K}_{anS}|$ is generally larger than $|\tilde{K}_{an1}|$, reflecting the enhanced symmetry-breaking at the nanodot boundary. This two-constant formulation is essential for accurately capturing the finite-size and surface-dominated effects that critically govern the stability of emergent topological textures in confined nanoscale geometries $[?]$.

Although the parameter vector θ governs the full three-dimensional spin field, the observable exploited in this work is restricted to its out-of-plane projection, the s_z component. This choice is dictated by experimental measurability: the most widely available magnetic imaging modalities, such as magnetic force microscopy and the polar magneto-optical Kerr effect, are primarily sensitive to the out-of-plane magnetization $[?]$, whereas resolving the in-plane components requires substantially more specialized techniques. Working on s_z therefore keeps the framework consistent with the information realistically accessible from experiment. This restriction, however, imposes a theoretical limit on parameter identifiability: distinct parameter sets, most notably those related by the sign of the antisymmetric exchange, which reverses the in-plane chirality while leaving s_z invariant, become indistinguishable in the observation space, so that a regressor exposed to both would collapse its predictions toward a degenerate average. To avoid this pathology, the present study deliberately samples the Hamiltonian parameters over value ranges that remain well-posed under the s_z -only observation, ensuring that the retained configurations do not induce the aforementioned averaging saturation; the specific bounds and the associated sign convention are detailed in Section ??.

3.2. GPU-Accelerated Image Dataset Generation

To capture the physical thermodynamic behavior and algorithmically drive the system toward a low-energy, thermodynamically stable configuration of the nanodot, thermal fluctuations must be formally introduced through the framework of statistical mechanics. While the atomistic Hamiltonian $\mathcal{H}$ rigorously defines the internal magnetic energy for a given spatial configuration of spins, it evaluates the system strictly as a frozen, zero-temperature snapshot. Hence, temperature is a macroscopic

parameter that governs the statistical distribution of the discrete magnetic microstates rather than acting as a localized interaction within the Hamiltonian. Consequently, temperature T [K] couples to the system via the Boltzmann distribution [?]. Formally a magnetic microstate μ represents a specific, instantaneous realization of the continuous spin matrix, denoted as $\mathbf{S}_\mu$. The total thermodynamic energy of this specific configuration is evaluated against the static environmental vectors and physical interaction strengths, such that $\mathcal{H}(\mu) \equiv \mathcal{H}(\mathbf{S}_\mu | \hat{\mathbf{d}}_{ij}, \hat{\mathbf{h}}, \mathbf{E}; \theta)$. At thermal equilibrium, the probability $P(\mu)$ of the system occupying this specific microstate is defined as:

$$P(\mu) = \frac{1}{\mathcal{Z}} \exp\left(-\frac{\mathcal{H}(\mu)}{k_B T}\right), \quad (5)$$

where k_B is the Boltzmann constant (expressed in meV/K). The term $\mathcal{Z}$ represents the canonical partition function, which mathematically enforces probability normalization by summing over all valid spin configurations ν (i.e., all possible state matrices $\mathbf{S}_\nu$ strictly satisfying the geometric boundary constraint $\|\mathbf{s}_i\|_2 = 1$).

$$\mathcal{Z} = \sum_{\mathbf{S}_\nu} \exp\left(-\frac{\mathcal{H}(\mathbf{S}_\nu | \hat{\mathbf{d}}_{ij}, \hat{\mathbf{h}}, \mathbf{E}; \theta)}{k_B T}\right). \quad (6)$$

In our computational framework, this thermodynamic distribution is actively sampled using the Metropolis-Hastings Monte Carlo algorithm combined with a simulated annealing protocol [?]. This stochastic sampling ensures a rigorous exploration of the highly multimodal energy landscape characteristic of these degenerate magnetic systems. When a localized spin perturbation proposes a physical transition from the current microstate $\mu^{(t)}$ (characterized by $\mathbf{S}_{\mu^{(t)}}$) to a newly modified state $\mu^{(t+1)}$ (characterized by $\mathbf{S}_{\mu^{(t+1)}}$), the resulting change in the system's total energy is computed as the scalar difference, $\Delta\mathcal{H} = \mathcal{H}(\mu^{(t+1)}) - \mathcal{H}(\mu^{(t)})$. To optimize computational efficiency, $\Delta\mathcal{H}$ is evaluated locally at the perturbed lattice site. The proposed state is subsequently accepted or rejected based on the precise transition probability:

$$P(\mu^{(t)} \rightarrow \mu^{(t+1)}) = \begin{cases} 1 & \text{if } \Delta\mathcal{H} < 0 \\ \exp\left(-\frac{\Delta\mathcal{H}}{k_B T^{(t)}}\right) & \text{if } \Delta\mathcal{H} \geq 0. \end{cases} \quad (7)$$

Physically, the instantaneous temperature parameter $T^{(t)}$ actively scales the energetic barriers within the system's vast topological energy landscape during the annealing process [?]. At elevated temperatures, the available thermal energy $k_B T^{(t)}$ frequently exceeds local energy penalties ($\Delta\mathcal{H}$), allowing the simulation to accept thermodynamically unfavorable spin alignments. This thermal noise provides the critical kinetic mechanism required to break out of metastable, localized energy minima such as misaligned ferromagnetic domains [?]. As $T^{(t)}$ is systematically scaled down during the annealing procedure, the probability of accepting higher-energy configurations undergoes exponential decay. This controlled thermal relaxation drives the continuous spin vector field toward a low-energy, thermodynamically stable configuration consistent with the annealing schedule, promoting the stabilization of complex chiral skyrmion textures driven by the underlying DMI.

Moreover, to ensure that each simulation consistently originates from a fully randomized paramagnetic phase regardless of the specific parameter configuration, the initial annealing temperature (i.e., $T^{(t)}$ at $t = 0$) is established as an exchange-normalized initial temperature. Consistent with foundational statistical mechanics [?], this starting temperature is dynamically rescaled in direct proportion to the dominant symmetric exchange energy ($\tilde{J}$). This energetic normalization guarantees that the initial thermal fluctuations are strictly sufficient to overcome the highest localized energetic barriers ($\Delta\mathcal{H}$) for that specific Hamiltonian. This effectively erases initial state bias and prevents the continuous spin vector field from prematurely collapsing into metastable local minima before the controlled cooling protocol begins.

The stochastic sampling above is realized as a fully vectorized, GPU-resident computation, one of the core methodological contributions of this work. Leveraging the matrix formulation of Section ??, the local energy variation $\Delta\mathcal{H}$ is evaluated not site by site, as in conventional simulators, but for the entire lattice at once, gathering neighbor contributions through tensor-shift operations combined as batched matrix products, and the full annealing trajectory is just-in-time compiled and executed on-device to eliminate host-device transfers. As illustrated in the inset of Figure ??, the parallelization operates on three levels, each acting only on independent or conditionally-independent degrees of freedom:

- *Spins*: Within a lattice, the nearest-neighbor Hamiltonian makes the two checkerboard sublattices A/B (defined by $(x + y + z) \bmod 2$) conditionally independent: all sites of one sublattice are updated in parallel with the other held fixed, before switching. This is equivalent to a sequential sweep, so it does not alter the transition $\mu^{(t)} \rightarrow \mu^{(t+1)}$.
- *Replicas*: The same parameter set θ is propagated with different random seeds as independent Markov chains, batched as an additional dimension on a GPU, providing robust thermodynamic averages.
- *Experiments*. Different parameter sets θ , physically independent systems, are distributed across multiple GPUs under a single-program multiple-data (pmap) scheme.

Crucially, the temporal Monte Carlo sweeps and the annealing temperature schedule remain sequential, since each transition depends on the previous state; only independent chains and conditionally-independent sites are computed concurrently. This parallel formulation reduces the wall-clock cost of dataset generation by orders of magnitude relative to sequential atomistic codes, without any approximation to the underlying physics.

To generate a robust and diverse dataset for deep learning evaluation, the magnetic parameters governing the Hamiltonian were systematically varied across the simulations. Key variables included the exchange interaction constant ($\tilde{J}$), the magnetocrystalline anisotropy ($\tilde{K}$), the external Zeeman field ($\tilde{H}_{ex}$), and the DMI strength ($\tilde{K}_{DM}$), alongside variations in the base initial temperature ($T^{(0)}$). Subsequently, the transformation of these raw, converged 3D spin states into normalized 2D image arrays followed a strict procedural pipeline.




Figure 1: Flowchart of the GPU-accelerated image-based magnetic dataset generation. Foundational inputs (green), the Hamiltonian parameters θ , the unit-cell geometry, and the external directional vectors, drive the atomistic Metropolis Monte Carlo simulation with simulated annealing (blue), relaxing the system into a converged 3D topological state. The resulting vector field is then projected, rendered, and preprocessed (yellow) into the labelled dataset $\mathcal{D}$ mapping 2D s_z images to their generating parameters. The bottom inset details the three levels of GPU parallelization: (a) *spins*, within a single lattice, the nearest-neighbor Hamiltonian makes the two checkerboard sublattices A/B conditionally independent, so all sites of one sublattice are updated simultaneously (with the other held fixed) before switching, reproducing a sequential sweep; (b) *replicas*, the same parameter set θ is run with different random seeds as independent Markov chains, batched on a GPU; and (c) *experiments*: different parameter sets θ are distributed across GPUs via single-program multiple-data (pmap) execution. The annealing temperature and the Monte Carlo sweeps $\mu^{(t)} \rightarrow \mu^{(t+1)}$ remain strictly sequential, so the parallelization introduces no approximation to the underlying physics.

First, during geometry parsing, the spatial coordinates defining the nanodot geometry were established using a binary circular lattice mask, with the disk radius and thickness fixed in terms of Magnetic Unit Cell (MUC) dimensions. Then, during spin-state mapping, the topological structure was projected onto a two-dimensional representation by extracting the out-of-plane component, s_z , of the continuous spin vectors $\mathbf{s}_i$ at all valid lattice coordinates. Next, the image-rendering stage isolated each configuration and rasterized it onto a fixed $\tilde{H} \times \tilde{W}$ pixel grid. A diverging red–blue colormap was used to encode opposite spin polarities, thereby visually capturing the chiral winding from the nanodot core to its perimeter. Also, as a preprocessing step, the pixel intensities associated with the s_z values were normalized to the interval $[0, 1]$ to promote stable gradient propagation during deep-learning optimization. Because every component of the target vector θ is a scalar quantity invariant under the discrete rotational symmetries of the lattice, the dataset was stored in its raw form and augmented on the fly during training through random (90°) rotations. As proper rotations, these transformations increase the effective sample diversity while strictly preserving both the physical parameters and the chiral handedness imposed by the Dzyaloshinskii–Moriya interaction (DMI).

Finally, the simulation pipeline yields a comprehensive image-based dataset, formalized as $\mathcal{D} = \{\mathbf{X}_n, \theta_n\}_{n=1}^N$, which serves as the final collection of input-output pairs for the deep learning parameter estimation task. Each input tensor

$\mathbf{X}_n \in [0, 1]^{\tilde{H} \times \tilde{W}}$ encodes the two-dimensional spatial magnetization field (the normalized s_z projection), while the corresponding target vector $\theta_n \in \mathbb{R}^P$ stores the P -dimensional array of generating thermodynamic and Hamiltonian parameters (e.g., $\theta_n = [T_n^{(0)}, \tilde{J}_{1,n}, \tilde{J}_{2,n}, \tilde{J}_{3,n}, \tilde{J}_{4,n}, \tilde{K}_{DM,n}, \tilde{H}_{ex,n}, \tilde{K}_{an1,n}, \tilde{K}_{anS,n}]$). Figure ?? depicts the main flowchart of our image-based magnetic dataset generation approach.

3.3. Deep Regression Framework

The inverse estimation of the Hamiltonian parameters and thermodynamic state from magnetic domain images is formalized as a continuous mapping problem. Given the generated dataset $\mathcal{D}$, where $\mathbf{X}_n$ represents the normalized s_z spatial magnetization field and θ_n represents the vector of generating physical parameters, we aim to learn a parameterized non-linear mapping function $f_\theta : \mathbb{R}^{\tilde{H} \times \tilde{W}} \rightarrow \mathbb{R}^P$, to find $\hat{\theta}_n \in \mathbb{R}^P$.

Here, f_θ denotes a deep neural network architecture parameterized by a set of learnable weights and biases. The optimal parameter set $\hat{\theta}$ is obtained by minimizing a continuous loss function over the training distribution [?]:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{N} \sum_{n=1}^N \mathcal{L}(\theta_n, f_\theta(\mathbf{X}_n)), \quad (8)$$

where $\mathcal{L}$ represents the chosen regression objective, typically the Mean Squared Error (MSE) or Huber loss, which mitigates

the influence of outlier topological states. By minimizing this objective, the network intrinsically learns to untangle the highly multimodal spatial features mapping to the specific energetic balances of the extended Heisenberg Hamiltonian.

To achieve robust feature extraction across varying scales of magnetic textures (e.g., from localized domain walls to global skyrmion geometries), this framework evaluates a conventional convolutional baseline alongside four state-of-the-art deep learning architectures: ResNet50, DenseNet, Xception, and the Visual Transformer (ViT). All models are trained from scratch on the magnetic dataset, without any transfer learning from natural-image collections. The four convolutional backbones (the baseline CNN, ResNet50, DenseNet, and Xception) share a common pipeline, a convolutional feature extractor followed by global average pooling and the regression head, and differ only in the internal structure of their convolutional block, formalized in the equations below. The ViT departs from this scheme, replacing convolutions with self-attention over image patches.

CNN (baseline): To quantify the accuracy attainable without transfer learning or the specialized inductive biases of the architectures below, we include a compact convolutional network trained from scratch directly on the native-resolution s_z maps. It stacks four convolutional stages of increasing width (each comprising two 3×3 convolutions with batch normalization, ReLU activation, and spatial dropout, interleaved with spatial max-pooling), followed by global average pooling and the shared regression head. Lacking residual, dense, depthwise separable, or attention mechanisms, this baseline isolates the marginal contribution of those components to the inverse estimation and serves as the reference against which the deeper architectures are compared.

ResNet50 [?]: Standard deep convolutional networks often suffer from the vanishing gradient problem as depth increases, hindering the learning of complex physical mappings. ResNet50 circumvents this by introducing shortcut connections that bypass multiple convolutional layers, directly propagating spatial information and gradients. Mathematically, if $\mathbf{x}$ is the input to a sequence of layers, and $\mathcal{F}(\mathbf{x}, \{\mathbf{W}_l\})$ represents the residual mapping to be learned (comprising convolutions, batch normalization, and ReLU activations) for the l -th layer, the output $\mathbf{y}$ of the residual block is defined as:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{\mathbf{W}_l\}) + \mathbf{x}. \quad (9)$$

In the context of magnetic domain images, these skip connections allow the network to simultaneously preserve low-level spatial features (like abrupt chiral boundaries) while composing deep, high-level abstractions of the global thermodynamic state.

DenseNet [?]: While ResNet adds features via summation, DenseNet concatenates them, maximizing information flow and feature reuse throughout the network. This is particularly advantageous for the Hamiltonian inverse problem, as the delicate energetic competition often manifests in subtle, overlapping spatial frequencies that require dense multi-scale analysis. In a DenseNet block, the feature maps of all preceding layers are used as inputs to the current layer. If $\mathbf{x}_l$ denotes the output of the l -th layer, and $H_l(\cdot)$ represents a composite function of

batch normalization, ReLU, and convolution, the forward pass is expressed as:

$$\mathbf{x}_l = H_l([\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_{l-1}]), \quad (10)$$

where $[\dots]$ denotes the concatenation operation. This dense topological structure ensures that the final regression head has direct access to both fine-grained pixel interactions and coarse global structures.

Xception [?]: This architecture operates on the hypothesis that cross-channel correlations and spatial correlations in feature maps can be entirely decoupled. It replaces standard Inception modules with depthwise separable convolutions, drastically reducing computational overhead while enhancing representational efficiency. A depthwise separable convolution splits the standard convolution into two distinct steps. First, a depthwise convolution applies a single spatial filter to each input channel independently. Second, a pointwise convolution (1×1 convolution) projects the output of the depthwise step onto a new channel space to capture cross-channel correlations. If $\mathbf{X} \in \mathbb{R}^{\tilde{H} \times \tilde{W} \times \tilde{C}}$ is the input tensor, the operation is formally approximated as:

$$\text{Conv}_{\text{sep}}(\mathbf{X}) = \text{Conv}_{\text{pointwise}}(\text{Conv}_{\text{depthwise}}(\mathbf{X})). \quad (11)$$

For nanoscale magnetism, this extreme decoupling allows the network to efficiently isolate specific topological frequency patterns (depthwise) before correlating them to the complex parameter space θ_n .

ViT [?]: Unlike CNNs, which enforce a strict local inductive bias, the Vision Transformer processes the 2D spatial magnetization field $\mathbf{X}_n$ as a sequence of flattened, non-overlapping 2D patches. This allows the model to capture long-range, global dependencies across the nanodot instantly, which is critical for understanding macro-scale topological stability dictated by the Zeeman field or global anisotropy. The input image is divided into N_p patches, which are linearly projected into D -dimensional embeddings. The core mechanism is the Multi-Head Self-Attention (MSA), which computes attention scores to weigh the importance of different patches relative to one another. For a given set of queries Q , keys K , and values V derived from the patch embeddings, the scaled dot-product attention is defined as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V, \quad (12)$$

where d_k is the scaling factor (the dimension of the keys). By employing MSA, the ViT can correlate a localized structural defect on one edge of the nanodot with a compensatory chiral boundary on the opposite edge, mapping these distant correlations directly to the global DMI or symmetric exchange strengths.

To optimize the parameter set $\hat{\Theta}$ in Equation ?? across all evaluated architectures, the models are trained utilizing the MSE as the explicit continuous loss function $\mathcal{L}$. The optimization process is executed using the Adam optimizer, relying on back-propagation and automatic differentiation to efficiently compute the network gradients [?]. By iteratively updating the model weights based on these computed gradients, the training pipeline systematically minimizes the structural and energetic discrepancies between the predicted outputs and the true Hamiltonian

constants, ensuring robust convergence in estimating the underlying physical parameters. All architectures, the conventional baseline and the four deeper models, are trained from scratch under an identical protocol, including the on-the-fly 90° rotation augmentation of the input s_z maps, so that any observed performance difference reflects the architecture itself rather than pretraining or data-pipeline discrepancies.

3.4. Spatial Interpretability via RAMs

Interpretability methods based on activation mapping have become a foundational approach for understanding the decision-making processes of deep neural networks. By identifying the spatial regions within an input that most significantly influence a model's prediction, these methods project opaque latent representations back into a human-interpretable domain. Among these, Gradient-weighted Class Activation Mapping (Grad-CAM) and its variants have been extensively utilized in computer vision to generate class-specific saliency maps [?].

In standard classification paradigms, given an input image $\mathbf{X}$ and a discrete target class score y^c , activation maps are computed by extracting the gradients of y^c with respect to the k -th feature map $\mathbf{A}^k \in \mathbb{R}^{H' \times W'}$ of a terminal convolutional layer (where H' and W' denote the spatial dimensions of the feature map). The spatial importance weights are derived by globally pooling these gradients, and the final map is a ReLU-activated linear combination of the feature maps [?]. While highly effective for discrete semantic categorization, this standard formulation structurally fails in continuous regression tasks. The absence of a discrete target class removes the natural selection mechanism required to isolate the output dimension driving the interpretability process. Furthermore, directly computing gradients with respect to a prediction $\hat{\theta}$ introduces severe gradient instability; minute fluctuations in the scalar output can yield drastically different spatial attributions, failing to provide the physical consistency required for analyzing complex Hamiltonian parameters [?].

3.4.1. RAMs for CNN-based Approaches

To overcome the fundamental limitations of discrete activation methods in highly degenerate physical systems, we introduce RAMs, a novel formulation explicitly tailored to bridge continuous parameter estimation with spatial topologies. Given an input magnetic domain visualization $\mathbf{X}$ and the optimized regression model $f_{\hat{\theta}} : \mathbb{R}^{\tilde{H} \times \tilde{W}} \rightarrow \mathbb{R}^{\tilde{P}}$, the network predicts a vector of physical constants $\hat{\theta} = f_{\hat{\theta}}(\mathbf{X})$. Let $\hat{\theta}_p \in \hat{\theta}$ denote the continuous prediction for the p -th target Hamiltonian parameter (e.g., the DMI strength $\tilde{K}_{DM}$ or symmetric exchange $\tilde{J}$), and let $\theta_p \in \theta$ represents its corresponding ground truth. Instead of utilizing a raw predicted scalar, we formulate a scale-invariant, error-based activation objective O_p . To prevent the mathematical discontinuities and negative objective bounds inherent to absolute error formulations $(1 - |\theta_p - \hat{\theta}_p|)$, we utilize a normalized Gaussian decay function:

$$O_p = \exp\left(-\gamma_p (\theta_p - \hat{\theta}_p)^2\right), \quad (13)$$

where $\gamma_p > 0$ is a hyperparameter-specific precision factor determining the strictness of the error decay. This objective smoothly

bounds the activation target strictly within $(0, 1]$, promoting maximum gradient response precisely when the network's prediction asymptotically approaches the true thermodynamic value, thereby isolating the topological features that positively contribute to accurate regression.

The spatial importance weights $\alpha_k^{(p)}$ for each continuous feature map $\mathbf{A}^k$ are subsequently computed by globally averaging the gradients of the error objective O_p with respect to the local spatial activations $\mathbf{A}_{ij}^k$:

$$\alpha_k^{(p)} = \frac{1}{H'W'} \sum_{i=1}^{H'} \sum_{j=1}^{W'} \frac{\partial O_p}{\partial \mathbf{A}_{ij}^k}. \quad (14)$$

These weights strictly quantify the importance of the k -th latent spatial feature in successfully minimizing the prediction error for the target physical parameter θ_p . The resulting RAM is then formulated as:

$$\mathcal{R}_{\text{RAM}}^{(p)} = \text{ReLU}\left(\sum_k \alpha_k^{(p)} \mathbf{A}^k\right). \quad (15)$$

The application of the ReLU activation ensures that only spatial features exerting a positive, stabilizing influence on the estimation of the target parameter are preserved. Finally, $\mathcal{R}_{\text{RAM}}^{(p)}$ is spatially upsampled to the original input resolution $\tilde{H} \times \tilde{W}$ and normalized to the interval $[0, 1]$. Then, by mapping the continuous error gradients back to the input spatial magnetization field $\mathbf{X}$, RAMs allow for the direct visual and quantitative identification of the specific emergent textures that the network utilizes to deduce the underlying energetic balances.

3.4.2. RAMs for ViT-based Approaches

Extending our RAMs-based interpretability to the ViT requires mapping sequence-based latent representations back to the two-dimensional physical domain. Unlike CNNs, which preserve spatial locality through sliding convolutional filters, the ViT processes the input spatial magnetization field $\mathbf{X}$ as a flattened sequence of N_r non-overlapping patches, outputting a discrete sequence of latent embeddings. Then, we introduce a token-reshaping mechanism that bridges the one-dimensional transformer latent space with the two-dimensional topological geometry. Let the output of the terminal Transformer encoder block be denoted as $\mathbf{Z} \in \mathbb{R}^{(N_r+1) \times D}$, where D is the embedding dimension and the extra token corresponds to the global aggregation token (e.g., the [CLS] token) utilized by the regression head to predict the Hamiltonian parameters $\hat{\mathbf{y}}$. By isolating the N_r spatial tokens, we obtain the localized sequence $\mathbf{Z}_s \in \mathbb{R}^{N_r \times D}$. Because each sequence token explicitly corresponds to a fixed physical patch of the original nanodot domain, $\mathbf{Z}_s$ can be folded back into a 2D spatial grid. Assuming square patches, the sequence is reshaped into a pseudo-feature map:

$$\mathbf{A}_{\text{ViT}} = \text{Reshape}(\mathbf{Z}_s) \in \mathbb{R}^{\sqrt{N_r} \times \sqrt{N_r} \times D}. \quad (16)$$

Once projected into this tensor format, $\mathbf{A}_{\text{ViT}}$ functions structurally identically to the terminal convolutional feature map $\mathbf{A}$ utilized in the CNN baseline. The k -th slice of the embedding dimension ($k \in \{1, \dots, D\}$) forms the latent spatial map

820 $\mathbf{A}_{\text{ViT}}^k \in \mathbb{R}^{\sqrt{N_r} \times \sqrt{N_r}}$. We then apply the identical scale-invariant, 867
 821 error-based activation objective O_p defined in Equation ?? . The 868
 822 spatial importance weights for the ViT embeddings are com- 869
 823 puted by globally averaging the gradients of this objective with 870
 824 respect to the reshaped spatial tokens:

$$\alpha_k^{(p)} = \frac{1}{N_r} \sum_{i=1}^{\sqrt{N_r}} \sum_{j=1}^{\sqrt{N_r}} \frac{\partial O_p}{\partial (\mathbf{A}_{\text{ViT}}^k)_{ij}}. \quad (17)$$

825 RAMs for the ViT are constructed via the ReLU-activated 873
 826 linear combination of these reshaped embedding slices, pre- 874
 827 cisely mirroring Equation ?? , before being upsampled to the 875
 828 original input resolution $\tilde{H} \times \tilde{W}$. This extension guarantees an 876
 829 architecture-agnostic evaluation of interpretability, ensuring that 877
 830 the isolated degenerate textures driving the model’s predictions 878
 831 can be fairly compared across fundamentally disparate deep 879
 832 learning architectures. 880

833 3.4.3. Normalized RAMs for Layer-Wise Analysis

834 To rigorously compare the spatial activations across different 883
 835 Hamiltonian parameters and network depths, it is imperative 884
 836 to compute the RAMs utilizing strictly normalized predictions 885
 837 and targets. Let $\bar{\theta}_p \in [0, 1]$ and $\hat{\theta}_p \in [0, 1]$ denote the Min- 886
 838 Max normalized ground truth and predicted values for the p -th 887
 839 physical parameter, respectively. By scaling the variables to 888
 840 a uniform interval, the layer-specific error objective is strictly 889
 841 bounded and evaluated as $O_p = \exp(-\gamma_p(\bar{\theta}_p - \hat{\theta}_p)^2)$. For a given 890
 842 convolutional or transformer layer $l \in \{1, \dots, L\}$, we extract the 891
 843 layer-specific unnormalized RAM, denoted as $\mathcal{R}_{\text{RAM}}^{(l,p)}$, following 892
 844 the gradient and activation protocols established in Equations ?? 893
 845 and ?? . To quantitatively evaluate and isolate the relative contri- 894
 846 bution of each distinct layer to the estimation of every physical 895
 847 constant, we introduce a global cross-layer and cross-parameter 896
 848 normalization scheme. The globally normalized activation map, 897
 849 $\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)} \in [0, 1]$, is defined by scaling the layer-specific maps 898
 850 against the absolute maximum spatial activation observed across 899
 851 all L layers and all $\tilde{P}$ parameters:

$$\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)} = \frac{\mathcal{R}_{\text{RAM}}^{(l,p)}}{\max_{l' \in \{1, \dots, L\}, p' \in \{1, \dots, \tilde{P}\}} \left(\max_{i,j} [\mathcal{R}_{\text{RAM}}^{(l',p')}]_{i,j} \right)}. \quad (18)$$

852 This normalization ensures that gradient magnitudes are di- 905
 853 mensionally consistent, preventing parameters with higher vari- 906
 854 ance or layers with larger gradient flows from dominating the 907
 855 interpretability metrics. Tracking $\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)}$ across network depth 908
 856 then reveals how the architecture resolves the ill-conditioned 909
 857 inverse problem, in which disparate combinations of $\tilde{J}$, $\tilde{K}_{DM}$, $\tilde{K}_{\text{S}}$ 910
 858 and $\tilde{H}_{ex}$ yield nearly degenerate states, through a hierarchical 911
 859 physical decoupling. Early layers (lower l) exhibit high activa- 912
 860 tions for short-range interactions, capturing the localized chiral 913
 861 winding and domain-wall boundaries dictated by the competi- 914
 862 tion between $\tilde{J}$ and $\tilde{K}_{DM}$, whereas deeper layers aggregate these 915
 863 into macro-scale representations, concentrating on the globally 916
 864 coupled parameters ($T^{(0)}$, exchange, and field terms) that govern 917
 865 the overall order and phase transitions. Correlating the intensity 918
 866 and localization of $\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)}$ across the hierarchy thus exploits both 919

local boundary chirality and global domain spacing to disam-
 biguate the Hamiltonian parameters.

The complete workflow for our RAMs-based localized spatial interpretability is illustrated in Figure ??.

871 4. Experimental Set-Up

To evaluate the effectiveness of the proposed framework for
 Hamiltonian parameter estimation, we design a comprehensive
 experimental protocol that integrates physics-based simulations,
 data-driven modeling, and interpretability analysis. We aim to
 assess not only the predictive performance of the model, but
 also its ability to generalize across different physical regimes
 and provide spatial explanations through the proposed RAMs.
 The dataset utilized in this study consists entirely of simulated
 magnetic domain configurations generated through our atomistic
 Monte Carlo framework, as described in Section ?? . To accu-
 rately reflect the finite-size effects and surface boundaries inher-
 ent to experimental nanostructures, the magnetic nanodots were
 modeled on a three-dimensional discrete lattice. The cylindrical
 geometry was constrained by a radius of $R_d = 18.25$ [MUC]
 and a thickness of $\tilde{L} = 5$ [MUC]. Within this confined grid, the
 spatial bounds were dictated by a strictly defined circular binary
 mask, yielding an active interacting volume of approximately
 5,200 localized spins per simulated dot.

882 4.1. Atomistic Simulations Details

For each unique combination of Hamiltonian parameters, the
 system was subjected to a rigorous simulated annealing protocol
 to isolate the thermodynamically relaxed state. To guarantee
 full ergodicity, the exchange-normalized initial temperature was
 established at an effective paramagnetic ceiling. The system was
 systematically cooled to a near-ground state of $T^{(f)} = 0.1$ K
 utilizing a linear decrement of $\Delta T = -0.1$ K. At each discrete
 temperature interval, the spin lattice underwent 10,000 Monte
 Carlo Sweeps (MCS) exclusively dedicated to thermalization,
 thereby overcoming metastable localized energetic barriers. This
 thermalization phase was subsequently followed by 2,000 mea-
 surement MCS to compute macroscopic thermodynamic aver-
 ages (e.g., specific heat and magnetic susceptibility) and verify
 strict energetic convergence. To overcome the prohibitive com-
 putational bottleneck traditionally associated with sequential
 Monte Carlo methods, the atomistic simulation was rigorously
 optimized implemented in JAX as a fully vectorized, just-in-
 time-compiled routine distributed across multiple GPUs via
 pmap parallelism in single precision.

Upon converging to a low-energy thermodynamically stable
 state at 0.1 K, the final topological state was isolated. To mimic
 experimental 2D observation techniques, the spatial magneti-
 zation field was extracted from the central discrete layer of
 the nanodot ($z = \lfloor \tilde{L}/2 \rfloor$). The out-of-plane spin components
 ($s_z \in [-1, 1]$) were normalized to the strict interval $[0, 1]$ and
 mapped onto a Cartesian pixel grid, finalizing the input im-
 ages $\mathbf{X}$ corresponding to the physical targets $\boldsymbol{\theta}$. The extended
 Hamiltonian parameter vector $\boldsymbol{\theta} \in \mathbb{R}^8$ estimated in this study is
 defined as: $\boldsymbol{\theta} = [T^{(0)}, \tilde{J}_2, \tilde{J}_3, \tilde{J}_4, \tilde{K}_{an1}, \tilde{K}_{anS}, \tilde{H}_{ex}, \tilde{K}_{DM}]$, where




Figure 2: RAMs extraction pipeline. Latent feature maps and continuous predictions from the network backbone are explicitly coupled via an error-decay objective. Gradient-based importance weights combine these features to produce a normalized, physically traceable spatial saliency map for the target Hamiltonian parameter θ_p .

the first-shell exchange constant $\tilde{J}_1 = 1.0$ meV is fixed as the energy reference scale to break the scale symmetry of the inverse problem and ensure a well-posed regression target (see Section ??). Each parameter was drawn with a quasi-random (low-discrepancy) sampling scheme within the physically motivated bounds summarized in Table ??, ensuring a systematic and homogeneous coverage of the admissible parameter region

Table 1: Hamiltonian parameters sampling ranges. The first-shell exchange $\tilde{J}_1 = 1.0$ meV is fixed as the energy reference. Parameters are drawn with a quasi-random (low-discrepancy) scheme within the specified bounds. Physical units: temperature in Kelvin [K], exchange and DMI constants in meV, anisotropy constants in meV/atom, external field in meV/atom.

Parameter	Physical role	Min	Max
$T^{(0)}$ [K]	Annealing temperature	0.0	20.0
$\tilde{J}_2$ [meV]	2nd-shell exchange	-0.33	0.66
$\tilde{J}_3$ [meV]	3rd-shell exchange	-0.29	0.29
$\tilde{J}_4$ [meV]	4th-shell exchange	-0.23	0.24
$\tilde{K}_{an1}$ [meV/atom]	Bulk anisotropy	0.0	4.57
$\tilde{K}_{anS}$ [meV/atom]	Surface anisotropy	0.0	0.20
$\tilde{H}_{ex}$ [meV/atom]	External Zeeman field	0.0	1.20
$\tilde{K}_{DM}$ [meV]	DMI strength	0.0	1.20

It is important to note that while all eight parameters are simultaneously estimated by the regression models, their effective identifiability from the s_z projection varies substantially due to the physical constraints of the observational modality. Six of the eight parameters, the annealing temperature $T^{(0)}$, the second-shell exchange $\tilde{J}_2$, the DMI strength $\tilde{K}_{DM}$, the bulk and surface anisotropy constants $\tilde{K}_{an1}$ and $\tilde{K}_{anS}$, and the external field $\tilde{H}_{ex}$, are effectively recovered from the two-dimensional projection. This is made possible by the deliberate enrichment of the dataset: applying the external field along the out-of-plane axis in part of the simulations couples the Zeeman term directly to the observable s_z component, while the extended anisotropy range produces sufficiently distinct domain-wall morphologies for the anisotropy constants to leave a measurable spatial signature. In contrast, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ remain essentially unidentifiable: their energetic contributions are strongly dominated by the first- and second-shell interactions $\tilde{J}_1$ and $\tilde{J}_2$, rendering their individual spatial signatures indistinguishable in the magnetization texture. These two constants are nevertheless retained in the regression target θ to provide a complete and physically honest characterization of the inverse problem, explicitly delineating its fundamental identifiability limits rather than concealing them. Their inclusion does not degrade the accuracy of the six recoverable parameters, since the

multi-output regression head learns to allocate gradient capacity independently to each output dimension. The resulting dataset is publicly available and can be accessed through the Kaggle repository¹. It contains a total of 166,141 samples, where each input is a grayscale magnetic domain image of size $39 \times 39 \times 1$ pixels, and each corresponding target is a vector of eight physical parameters $\theta \in \mathbb{R}^8$ as defined above.

To broaden the structural diversity of the dataset beyond a single field geometry, the external field orientation was varied within the dataset: in addition to the in-plane configuration, an out-of-plane field was employed to stabilize additional topologies such as isolated skyrmions, while the magnetocrystalline anisotropy range was deliberately extended toward high values to access an Ising-like regime characterized by sharper, more discrete domain walls. The dataset also spans two complementary mechanisms for stabilizing modulated, non-collinear textures: a chiral regime driven by the DMI $\tilde{K}_{DM}$ and an exchange-frustration regime driven by the second-shell exchange $\tilde{J}_2$. To isolate the contribution of each mechanism, these two interactions are activated in a mutually exclusive fashion across the sampled configurations. Consequently, in the pairwise correlation matrix of the eight parameters, $\tilde{J}_2$ and $\tilde{K}_{DM}$ are the only pair exhibiting an appreciable anti-correlation ($\rho \approx -0.53$), a direct consequence of this deliberate design, while all remaining pairs are essentially uncorrelated ($|\rho| \lesssim 0.3$), confirming that the model cannot exploit spurious cross-parameter correlations to infer one parameter from another. This controlled heterogeneity enriches the space of magnetic phases that the regression models must disentangle and, together with the sampling bounds of Table ??, exposes the models to a wide range of competing energetic regimes within a single unified dataset.

The parameter bounds are not intended to reproduce a specific compound but to cover the region of Hamiltonian space where chiral and modulated magnetic textures exist. Since $\tilde{J}_1$ is fixed as the energy reference, the physically meaningful quantities are the dimensionless ratios $\tilde{J}_r/\tilde{J}_1$, $\tilde{K}_{DM}/\tilde{J}_1$, and $\tilde{K}/\tilde{J}_1$, and these ratios fall within the ranges reported by first-principles calculations and experiments for chiral magnets and ultrathin transition-metal films. The higher-order exchange constants ($|\tilde{J}_2| > |\tilde{J}_3| > |\tilde{J}_4|$, with $\tilde{J}_2/\tilde{J}_1$ up to 0.66) reproduce the decaying, sign-alternating exchange obtained from density-functional calculations of ultrathin films such as Fe/Ir(111) and Pd/Fe/Ir(111) [?], and lie

¹<https://www.kaggle.com/datasets/carlosanamejoy/dataset-spines-complete>

in the strong-frustration regime of the classical $\tilde{J}_1$ - $\tilde{J}_2$ model. For the DMI, we sample $\tilde{K}_{DM}/\tilde{J}_1$ up to ≈ 1.2 ; while interfacial materials typically exhibit $\tilde{K}_{DM}/\tilde{J}_1 \sim 0.3$ - 0.7 [?], larger ratios are the standard choice in finite-lattice Monte Carlo studies so that the helical period ($\lambda \propto \tilde{J}_1/\tilde{K}_{DM}$) fits several times within the simulation cell [?]. Because the physics is scale-invariant, a large ratio on a small lattice is equivalent to a small ratio in a large system, so this choice reflects standard finite-size rescaling rather than unphysical parameter values. The anisotropy range (up to $\tilde{K}_{anl} \approx 4.6$ meV/atom) spans from soft magnets to the strongly anisotropic interfacial regime, consistent with values reported for systems ranging from FePt to single Co adatoms on Pt(111) [?]. Finally, the Zeeman range ($\tilde{H}_{ex}$ up to 1.2 meV/atom) corresponds, through $E = g\mu_B B$ with $g \approx 2$, to magnetic fields up to ≈ 10 T, the upper range of standard laboratory magnets used to map skyrmion phase diagrams, and the temperature range ($T^{(0)}$ up to 20 K, i.e., $k_B T/\tilde{J}_1$ up to ≈ 1.7) sweeps the system from the ordered ground state through the paramagnetic transition, consistent with ordering temperatures of chiral magnets with $\tilde{J}_1 \sim 1$ meV.

It is important to note that the DMI strength $\tilde{K}_{DM}$ is sampled exclusively over non-negative values ($\tilde{K}_{DM} \in [0.0, 1.2]$ meV). This restriction is a physically motivated convention arising from the symmetry properties of the antisymmetric exchange term in Equation ?? . Under the transformation $\tilde{K}_{DM} \rightarrow -\tilde{K}_{DM}$, the equilibrium spin configuration undergoes a reversal of its in-plane chirality, the rotational sense of the chiral winding inverts, while the out-of-plane magnetization component s_z at every lattice site remains strictly invariant. Since the regression models operate exclusively on the two-dimensional s_z projection, configurations with $+\tilde{K}_{DM}$ and $-\tilde{K}_{DM}$ are exactly degenerate in the observation space, rendering the sign of the DMI fundamentally unidentifiable from s_z images alone. Including both signs simultaneously would introduce a strict two-fold degeneracy in the input-output mapping, collapsing the model's predictions toward $\hat{K}_{DM} \approx 0$ and precluding meaningful regression. The chirality handedness is instead encoded through the fixed direction vector $\hat{\mathbf{d}}_{ij}$, consistent with standard conventions in atomistic modeling of interfacial DMI, where the sign of the antisymmetric exchange sets the preferred rotational sense of the chiral texture [?]. Recovering the DMI sign would require access to chirality-sensitive observables, such as the full three-dimensional spin vector field or in-plane magnetization components.

4.2. Deep Learning Training

To ensure robust evaluation and prevent information leakage during the deep learning optimization phase, the generated dataset $\mathcal{D}$ is partitioned into training, validation, and testing subsets using a rigorous two-stage random split strategy. Initially, 15% of the total dataset is isolated as a hold-out test set to evaluate ultimate model generalization. Subsequently, the remaining 85% is further divided, allocating exactly 15% of the total initial data for validation. This procedure yields an effective split of 70% for training, 15% for validation, and 15% for testing. All partitions are executed utilizing a fixed random seed to guarantee strict reproducibility across all experimental

runs. This strategy ensures that topological samples remain strictly independent across subsets while perfectly preserving the underlying thermodynamic distributions of the dataset.

Prior to network ingestion, both the input spatial magnetization images and the target Hamiltonian variables are preprocessed to ensure numerical stability and efficient gradient convergence. The physical target variables are scaled using Min-Max normalization, mapping each continuous parameter to the strict interval $[0, 1]$ based exclusively on the statistical distribution of the training set. This identical transformation is subsequently applied to the validation and test sets to maintain inferential consistency and directly support the bounded RAMs objective $O_p \in (0, 1]$. For the four deep backbone architectures, the single-channel s_z matrices are spatially resized to 224×224 pixels and duplicated across three channels (grayscale-to-RGB conversion). This dimensional expansion matches the input dimensionality expected by the standard deep backbone architectures; the conventional CNN baseline instead operates directly on the native $39 \times 39 \times 1$ maps. No pretrained weights are used, and all architectures are trained entirely from scratch.

The configuration of the evaluated deep learning architectures (see Section ??) is summarized in Table ?? . All five architectures share the same extended regression head designed to mitigate overfitting in the high-dimensional parameter space. Specifically, the global pooling output is passed through a Batch Normalization layer, followed by a Dropout layer with rate $p = 0.4$, a fully connected Dense layer of 256 units with ReLU activation and ℓ_2 weight regularization ($\lambda = 10^{-4}$), a second Batch Normalization layer, a second Dropout layer with rate $p = 0.3$, and a final linear Dense layer outputting the $\tilde{P} = 8$ target parameters. For the convolutional architectures (the conventional CNN, ResNet50, DenseNet121, and Xception), the global aggregation is performed via Global Average Pooling over the spatial feature maps, whereas for the ViT-B/16 the token sequence output of the terminal encoder block is first collapsed via Global Average Pooling along the sequence dimension before entering the shared regularization head. This design choice reflects the distinct spatial inductive biases of convolutional and attention-based backbones while ensuring a strictly comparable regression capacity across all evaluated models.

The optimization of the parameter set $\hat{\Theta}$ was executed utilizing the Adam optimizer with a unified initial learning rate of $\eta = 10^{-4}$ for all architectures, including the ViT-B/16, chosen for stable convergence when training all architectures from scratch. All models were trained to minimize the MSE loss function, while the MAE was actively monitored as a secondary diagnostic. Training was conducted using a highly parallelized batch size of 256 samples over a maximum of 100 epochs. To actively prevent overfitting and manage the learning dynamics within the highly degenerate optimization landscape, two primary callbacks were implemented: a dynamic learning rate scheduler (ReduceLROnPlateau) that reduced the learning rate by a factor of 0.3 if the validation loss stagnated for 5 consecutive epochs, and an EarlyStopping protocol that halted training and restored the best network weights if no validation improvement was observed after 12 consecutive epochs.

The predictive performance on the unseen test set was as-

Table 2: Configuration of evaluated deep learning architectures for inverse Hamiltonian parameter estimation. All models share a standardized regularized regression head mapped to the $\tilde{P} = 8$ -dimensional physical parameter space. BN: Batch Normalization; Drop(p): Dropout with rate p .

Model	Architecture Type	Pretraining	Global Aggregation	Output Head
Conventional CNN	Custom CNN	None (from scratch)	Global Avg. Pooling (spatial)	BN $\rightarrow$ Drop(0.4) $\rightarrow$ Dense(256, ReLU, ℓ_2) $\rightarrow$ BN $\rightarrow$ Drop(0.3) $\rightarrow$ Dense(8)
ResNet50	Residual CNN	None (from scratch)	Global Avg. Pooling (spatial)	BN $\rightarrow$ Drop(0.4) $\rightarrow$ Dense(256, ReLU, ℓ_2) $\rightarrow$ BN $\rightarrow$ Drop(0.3) $\rightarrow$ Dense(8)
DenseNet121	Densely Connected CNN	None (from scratch)	Global Avg. Pooling (spatial)	BN $\rightarrow$ Drop(0.4) $\rightarrow$ Dense(256, ReLU, ℓ_2) $\rightarrow$ BN $\rightarrow$ Drop(0.3) $\rightarrow$ Dense(8)
Xception	Depthwise Separable CNN	None (from scratch)	Global Avg. Pooling (spatial)	BN $\rightarrow$ Drop(0.4) $\rightarrow$ Dense(256, ReLU, ℓ_2) $\rightarrow$ BN $\rightarrow$ Drop(0.3) $\rightarrow$ Dense(8)
ViT-B/16	Patch-based Transformer	None (from scratch)	Global Avg. Pooling (tokens)	BN $\rightarrow$ Drop(0.4) $\rightarrow$ Dense(256, ReLU, ℓ_2) $\rightarrow$ BN $\rightarrow$ Drop(0.3) $\rightarrow$ Dense(8)

sessed using two complementary regression metrics. The coefficient of determination (R^2) quantifies the proportion of variance in each physical parameter successfully captured by the network:

$$R^2 = 1 - \frac{\sum_{n=1}^{N_{\text{test}}} (\theta_{n,p} - \hat{\theta}_{n,p})^2}{\sum_{n=1}^{N_{\text{test}}} (\theta_{n,p} - \bar{\theta}_p)^2}, \quad (19)$$

where $\bar{\theta}_p$ is the mean of the observed ground truth target. Note that R^2 is statistically undefined when the target variable exhibits near-zero variance within a given evaluation subset; in these cases the metric is reported as N/A. The MAE evaluates the absolute magnitude of the prediction error in the normalized feature space $[0, 1]$, providing a direct and interpretable measure of per-parameter accuracy:

$$\text{MAE} = \frac{1}{N_{\text{test}}} \sum_{n=1}^{N_{\text{test}}} |\theta_{n,p} - \hat{\theta}_{n,p}|. \quad (20)$$

Since all target variables are Min-Max normalized to $[0, 1]$ prior to training, the MAE is inherently dimensionless and directly comparable across all eight Hamiltonian parameters, irrespective of their original physical units or numerical ranges.

4.3. Spatial Interpretability: Implementation Details

The RAMs were computed on the held-out test set using the formulations introduced in Section ???. For the Xception backbone, RAMs were extracted from eight layers selected to systematically cover the three functional stages of the architecture: two entry-flow layers capturing low-level spatial features at high resolution, block1_conv2 (112×112) and block3_sepconv2 (56×56); one transitional layer at block4_sepconv2 (28×28); five middle-flow layers operating at the bottleneck resolution of 14×14 , block6_sepconv2, block8_sepconv2, block10_sepconv2, and block12_sepconv2; and one exit-flow layer, block13_sepconv2 (14×14), immediately preceding the global pooling operation. This stratified selection was motivated by the physical hypothesis that early layers encode fine-grained domain periodicity and local spin-winding patterns at the scale of individual domain walls, while middle-flow layers progressively aggregate these local features into intermediate topological representations, and exit-flow layers encode high-level semantic information about the global magnetic phase. By distributing the selected layers across all three functional stages, the analysis directly tests whether the hierarchical decoupling predicted by the normalized RAM formulation of Equation ??? is empirically realized in the Xception backbone. For the ViT, a single attribution map per parameter was derived from the spatial token embeddings at the output of the terminal encoder

block, reshaped into $\mathbf{A}_{\text{ViT}} \in \mathbb{R}^{14 \times 14 \times 768}$ following Equation ???. Since `keras_hub.ViTBackbone` does not expose internal token tensors as differentiable nodes within `tf.GradientTape`, the token sequence $\mathbf{Z}_s \in \mathbb{R}^{196 \times 768}$ was first extracted via a standard forward pass and then re-instantiated as a `tf.Variable` to enable gradient computation. The [CLS] token was treated as a fixed constant and the reconstructed sequence was passed through the `GlobalAveragePooling1D` and Dense regression head to compute O_p and its gradients with respect to $\mathbf{Z}_s$, following Equation ??.

The precision factor was fixed to $\gamma_p = 10$ uniformly for all $p \in \{1, \dots, 8\}$. Given that all target variables are Min-Max normalized to $[0, 1]$, this value ensures a smooth and well-conditioned gradient signal across the full prediction error range: $O_p \geq e^{-10} \approx 4.5 \times 10^{-5}$ at maximum error and $O_p = 1$ at perfect prediction. Lower values of γ_p would yield spatially diffuse attributions insensitive to prediction quality, while higher values would cause gradient saturation for the less identifiable parameters where errors are systematically larger. The chosen value maintains sensitivity across all eight parameters, from the most identifiable ($T^{(0)}$, $\tilde{J}_2$, and $\tilde{K}_{DM}$) to the physically degenerate ones ($\tilde{K}_{an1}$, $\tilde{K}_{anS}$, $\tilde{H}_{ex}$).

To confine activations to the physical extent of the nanodot, a circular binary mask $\mathcal{M} \in \{0, 1\}^{224 \times 224}$, defined as a disk of radius $r = 0.94 \times 112$ pixels centered at the image center, was applied to each RAM before normalization, setting all activations outside the disk to zero. A concentration score $\xi \in [0, 1]$, measuring the fraction of total activation energy within the disk, was used to select the most physically representative sample per magnetic phase: among the ten test images closest to the UMAP centroid of each phase cluster, the sample maximizing $\bar{\xi}$ averaged over all parameters and layers was chosen as the phase representative $\mathbf{X}_n^*$.

To rigorously assess the RAMs formulation, we benchmark it against three model-agnostic attribution methods adapted here to the continuous regression setting: the Local Interpretable Model-agnostic Explanations (LIME) [?], the Shapley Additive Explanations (SHAP) in their kernel form over superpixels [?], and a spatial-frequency Shapley estimator (SF-SHAP) that evaluates exact Shapley values over a compact set of radial frequency bands [?]. For a target parameter p , let $f_p : \mathbb{R}^{\tilde{H} \times \tilde{W}} \rightarrow \mathbb{R}$ denote the scalar map returning the p -th predicted parameter, $f_p(\mathbf{X}) = \hat{\theta}_p$. The input image is partitioned into d disjoint interpretable regions (superpixels) $\{\mathcal{X}_1, \dots, \mathcal{X}_d\}$ obtained by spatial segmentation. A coalition vector $\mathbf{z} \in \{0, 1\}^d$ encodes which regions are retained ($z_i = 1$) or masked to a reference baseline b ($z_i = 0$), and the map $h_{\mathbf{X}}(\mathbf{z})$ reconstructs the corresponding perturbed image.

LIME. This method fits a sparse linear surrogate $g(\mathbf{z}) = \phi_0 +$

1193 $\sum_{i=1}^d \phi_i z_i$ that locally approximates f_p in the neighborhood of $\mathbf{X}$ 1235
 1194 by minimizing a proximity-weighted least-squares objective: 1236

$$\phi^{\text{LIME}} = \arg \min_g \sum_z \pi_{\mathbf{X}}(\mathbf{z}) (f_p(h_{\mathbf{X}}(\mathbf{z})) - g(\mathbf{z}))^2 + \Omega(g), \quad (21)$$

1195 where $\pi_{\mathbf{X}}(\mathbf{z})$ is a proximity kernel that weights each perturbation 1240
 1196 by its similarity to the original image, and $\Omega(g)$ penalizes surro-
 1197 gate complexity to enforce sparsity. The coefficient ϕ_i is taken 1241
 1198 as the attribution of region i . 1242

1199 **SHAP.** Grounded in cooperative game theory, the Shapley 1242
 1200 value attributes to each region i the average marginal contribution 1243
 1201 it makes to the prediction across all possible coalitions: 1244

$$\phi_i = \sum_{S \subseteq \mathcal{I} \setminus \{i\}} \frac{|\mathcal{S}|! (d - |\mathcal{S}| - 1)!}{d!} (f_p(h_{\mathbf{X}}(\mathbf{z}_{S \cup \{i\}})) - f_p(h_{\mathbf{X}}(\mathbf{z}_S))), \quad (22)$$

1202 where $\mathcal{I} = \{1, \dots, d\}$ and $\mathbf{z}_S$ is the coalition activating only the 1250
 1203 regions in S . This attribution is the unique one satisfying local 1251
 1204 accuracy, missingness, and consistency, but its exact evaluation 1252
 1205 requires 2^d model queries. Since d equals the number of super- 1253
 1206 pixels, exact enumeration is intractable, and SHAP therefore 1254
 1207 estimates the Shapley values through the kernel formulation: 1255
 1208 a weighted linear regression over sampled coalitions using the 1256
 1209 SHAP kernel: 1257

$$\pi_{\text{SHAP}}(\mathbf{z}) = \frac{d - 1}{\binom{d}{|\mathbf{z}|} |\mathbf{z}| (d - |\mathbf{z}|)}, \quad (23)$$

1210 so that $\phi^{\text{SHAP}} = \arg \min_{\phi} \sum_{\mathbf{z}} \pi_{\text{SHAP}}(\mathbf{z}) (f_p(h_{\mathbf{X}}(\mathbf{z})) - \phi_0 - \sum_i \phi_i z_i)^2$; 1261
 1211 the resulting coefficients converge to the Shapley values of Equa- 1261
 1212 tion ?? while sampling only a subset of coalitions. 1262

1213 **SF-SHAP.** Rather than segmenting the image into superpixels 1262
 1214 this variant partitions it into a small number Q of concentric 1263
 1215 radial frequency bands, matched to the characteristic spatial 1264
 1216 scales of the magnetic texture. Because Q is small ($Q = 7$ here), 1265
 1217 the full coalition space of $2^Q = 128$ subsets can be enumerated 1266
 1218 exactly, so SF-SHAP evaluates the Shapley values of Equa- 1267
 1219 tion ?? without sampling approximation, at the cost of a coarser 1268
 1220 frequency-resolved (rather than pixel-resolved) attribution. 1269

1221 All three baselines are perturbation-based: each explanation 1270
 1222 demands hundreds to thousands of forward evaluations of f_p to 1271
 1223 probe the coalition space, whereas a RAM is obtained from a 1272
 1224 single backward pass through the network. 1273

1225 To quantitatively and fairly compare the spatial faithfulness of 1274
 1226 every attribution method, RAMs and the perturbation baselines 1275
 1227 alike, we adopt a deletion protocol based on the Most Relevant 1276
 1228 First (MoRF) criterion [?]. Given the relevance map produced 1277
 1229 by a method m for the parameter p , the $\tilde{H}\tilde{W}$ input locations are 1278
 1230 ranked in order of decreasing relevance, yielding the ordering 1279
 1231 $\pi_1^{(m)}, \pi_2^{(m)}, \dots, \pi_{\tilde{H}\tilde{W}}^{(m)}$. For a deletion fraction $\rho \in [0, \rho_{\max}]$, the 1280
 1232 $[\rho \tilde{H}\tilde{W}]$ most relevant locations are replaced by the reference 1281
 1233 baseline b , producing the perturbed image $\mathbf{X}^{(m)}(\rho)$, and the in- 1282
 1234 duced deviation of the prediction is measured as: 1283

$$\Delta_p^{(m)}(\rho) = |f_p(\mathbf{X}^{(m)}(\rho)) - f_p(\mathbf{X})|. \quad (24)$$

A faithful attribution concentrates relevance on the locations 1235
 that genuinely drive the prediction, so that removing them in- 1236
 duces a large and rapid deviation. The overall faithfulness of 1237
 method m for parameter p is therefore summarized by the area 1238
 under the deletion curve: 1239

$$\text{AUC}_p^{(m)} = \int_0^{\rho_{\max}} \Delta_p^{(m)}(\rho) d\rho, \quad (25)$$

where a larger $\text{AUC}_p^{(m)}$ indicates a more spatially localized and 1240
 physically faithful explanation; a random-relevance map is in- 1241
 cluded as a lower-bound control. This deletion score, together 1242
 with the per-explanation computational cost, provides the quan- 1243
 titative basis for the interpretability comparison. 1244

All experiments were conducted in a cloud environment 1245
 on RunPod. The computational setup was equipped with 1246
 an NVIDIA H200 GPU. The training pipeline was imple- 1247
 mented in Python using TensorFlow, leveraging GPU accel- 1248
 eration for efficient model training and inference. The atom- 1249
 istic state simulations were executed separately on the multi- 1250
 GPU JAX configuration described in Section ?? . The complete 1251
 source code is publicly available at <https://github.com/deathperminut/PaperInverseProblemEstimation>, orga- 1252
 nized into three modules: /simulation containing the atom- 1253
 istic Monte Carlo Hamiltonian simulation scripts implemented 1254
 in JAX with multi-GPU (pmap) parallelization; /regression 1255
 containing the training, validation, and evaluation pipelines for 1256
 all five deep learning architectures; and /rams containing the 1257
 Regression Activation Maps computation and visualization code 1258
 for both the Xception CNN and the Vision Transformer. 1259
 1260

5. Results and Discussion

5.1. Thermodynamic Simulation of Topologies

We present the main results of the proposed atomistic Hamil- 1261
 tonian simulation framework (see Section ??). The simulated 1262
 dataset spans the eight-dimensional parameter space θ over the 1263
 bounds summarized in Table ?? , sampled with deliberately het- 1264
 erogeneous densities: the higher-order exchange constants $\tilde{J}_2$, 1265
 $\tilde{J}_3$, and $\tilde{J}_4$ concentrate near zero, whereas the DMI strength $\tilde{K}_{DM}$ 1266
 is broadly populated across its full $[0, 1.2]$ meV range, and the 1267
 bulk anisotropy $\tilde{K}_{an1}$ and external field $\tilde{H}_{ex}$ exhibit heavy tails 1268
 toward high values that deliberately incorporate high-anisotropy, 1269
 Ising-like regimes and out-of-plane field configurations. These 1270
 bounded ranges and sampling heterogeneities explicitly govern 1271
 the topological diversity of the generated magnetic phase 1272
 space and, in turn, the learning dynamics, optimization stability, 1273
 and predictive limits of the downstream deep regression frame- 1274
 works when navigating the challenging regimes of the extended 1275
 Heisenberg Hamiltonian. 1276

To validate the physical consistency of the generated dataset, 1277
 Figure ?? illustrates the thermodynamic evolution of a repre- 1278
 sentative nanodot during Metropolis–Hastings Monte Carlo an- 1279
 nealing, under a fixed configuration ($\tilde{J}_1 = 1.0$, $\tilde{K}_{an1} = 0.10$, 1280
 $\tilde{K}_{DM} = 1.00$ meV, remaining terms zeroed; Table ??) that iso- 1281
 lates the dominant competition between the first-shell exchange 1282
 $\tilde{J}_1$ and the DMI $\tilde{K}_{DM}$. The main panel tracks the total energy 1283
 1284

$\mathcal{H}$ [meV] versus the annealing temperature $T^{(0)}$ [K], with insets showing the s_z field of the central layer at each state: the system evolves from a disordered paramagnetic texture at high $T^{(0)}$, through nucleating short-range domains, to a well-defined labyrinthine (helical) configuration at $T^{(0)} = 0.1$ K. The monotonic decrease of $\mathcal{H}$ confirms the thermodynamic relaxation toward a low-energy stable configuration, validating the Monte Carlo pipeline and the quality of the generated dataset.




Figure 3: Thermodynamic convergence of a representative nanodot during simulated annealing. The main panel shows the total Hamiltonian energy $\mathcal{H}$ [meV] versus the annealing temperature $T^{(0)}$ [K]; each inset shows the s_z field of the central layer at that state, tracing the transition from a disordered paramagnetic phase at high $T^{(0)}$ to a relaxed labyrinthine configuration at $T^{(0)} = 0.1$ K.

To characterize the structural diversity of the dataset, Figure ?? presents a two-dimensional phase diagram of the magnetization as a function of the annealing temperature $T^{(0)}$ and the DMI strength $\tilde{K}_{DM}$. In the low-temperature regime, the system stabilizes into well-defined ordered textures: increasing $\tilde{K}_{DM}$ drives the sequence from nearly uniform ferromagnetic states, through chiral stripe and labyrinthine domains, to isolated skyrmionic textures. These ordered configurations remain visually distinguishable, which is precisely what allows most Hamiltonian parameters to be recovered by the inverse models. The principal source of ill-conditioning is instead thermal: as $T^{(0)}$ approaches the effective Curie point T_C , thermal fluctuations overwhelm the interactions and the system collapses into a paramagnetic, spatially uncorrelated texture carrying little discriminative information. This temperature-driven loss of structure, rather than any intrinsic parameter degeneracy, constitutes the dominant challenge the dataset poses to the inverse estimation.

5.2. Deep Learning Regression Results

The predictive performance of the five evaluated architectures (the conventional CNN baseline, DenseNet, ResNet, Xception, and ViT) was assessed on the held-out test set comprising 15% of the total dataset. Each model was tasked with simultaneously inferring the complete extended parameter vector θ directly from the two-dimensional out-of-plane magnetization image $\mathbf{X}_n$, without any auxiliary physical information. As shown in Table ??, all five architectures perform competitively, with the differences




Figure 4: Two-dimensional phase diagram of the magnetization texture $\mathbf{X}_n$ as a function of the annealing temperature $T^{(0)}$ and the DMI strength $\tilde{K}_{DM}$. Each cell shows an enlarged representative image on a coarse sampling grid; the background color encodes $\tilde{K}_{DM}$.

on the identifiable parameters typically confined to a few percent of R^2 . Among them, Xception emerges as the strongest overall model: it attains the highest R^2 for six of the eight parameters, including the most identifiable ones, $T^{(0)}$, $\tilde{J}_2$, $\tilde{K}_{an1}$, and $\tilde{K}_{DM}$, while the ViT marginally leads on the external field $\tilde{H}_{ex}$ and the conventional CNN baseline on the surface anisotropy $\tilde{K}_{anS}$. The same ranking is largely mirrored in the MAE panels, where Xception yields the lowest absolute error for the physically dominant parameters $T^{(0)}$, $\tilde{K}_{an1}$, and $\tilde{K}_{DM}$. A particularly informative observation is that the conventional CNN baseline, trained from scratch and devoid of pretraining or specialized inductive biases, matches, and for some parameters even exceeds, the deeper backbones. This confirms that neither transfer learning from natural images nor the sophisticated residual, dense, depthwise, or attention mechanisms provide a decisive advantage for this task; the estimation accuracy is instead governed by the physical information that the s_z projection makes available. Physically, Xception's slight edge stems from its depthwise separable convolution mechanism, which efficiently decouples the extraction of spatial spin-texture frequencies from the cross-channel correlations required to map those features onto the Hamiltonian parameter space, yielding the most accurate point predictions for the dominantly identifiable parameters. The ViT, in turn, remains competitive by processing $\mathbf{X}_n$ as a sequence of non-overlapping spatial patches and computing multi-head self-attention across all patch pairs, explicitly capturing the long-range correlations, such as the global chirality imposed by $\tilde{K}_{DM}$ and the large-scale ordering driven by $\tilde{J}_1$, that are inaccessible to purely local convolutional receptive fields.

Figure ?? plots the mean R^2 over the identifiable parameters against the number of trainable weights of each architecture, spanning nearly two orders of magnitude in model size, from the compact conventional CNN (~ 1 M parameters) to the ViT

(~ 86 M parameters). The resulting trend is essentially flat: the conventional CNN already attains a mean R^2 close to that of the largest models, and increasing the parameter count by two orders of magnitude yields only a marginal improvement, with Xception offering the best accuracy–size trade-off. This saturation demonstrates that predictive accuracy is not bottlenecked by model capacity but by the physical information accessible in the s_z projection: once an architecture is expressive enough to read the relevant spatial features, additional parameters and more elaborate inductive biases provide diminishing returns, confirming that the recovery of the identifiable parameters is a robust, architecture-independent property of the data.




Figure 5: Mean R^2 over the identifiable Hamiltonian parameters versus the number of trainable weights (log scale) for the five architectures. The nearly flat trend, the compact CNN already matching models two orders of magnitude larger, indicates that accuracy is limited by the physical information in the s_z projection, not by model capacity.

Table ?? reveals a pronounced and physically interpretable heterogeneity in predictive accuracy across the parameter space. The initial annealing temperature $T^{(0)}$, the second-shell exchange $\tilde{J}_2$, the DMI strength $\tilde{K}_{DM}$, the bulk anisotropy $\tilde{K}_{an1}$, and the external field $\tilde{H}_{ex}$ consistently yield high R^2 scores across all architectures, reflecting the robust and spatially distinctive topological imprints they leave on the observable magnetization texture: $T^{(0)}$ governs the global degree of spin order, $\tilde{J}_2$ controls the second-neighbor exchange pathway that modulates the domain-scale collinearity, and $\tilde{K}_{DM}$ directly dictates the characteristic wavelength and chirality of the domain patterns. However, a seemingly contradictory observation emerges when examining the MAE rows in Table ??: despite their high R^2 scores, $T^{(0)}$, $\tilde{J}_2$, and $\tilde{K}_{DM}$ also exhibit comparatively larger absolute errors in their respective physical units, reaching up to approximately 0.49 K for $T^{(0)}$, 0.021 meV for $\tilde{J}_2$, and 0.037 meV for $\tilde{K}_{DM}$ in the worst-performing architecture. This apparent inconsistency is not a failure of the models but rather a direct consequence of the thermodynamic landscape of the dataset. As illustrated in the phase diagrams of Figure ??, the parameter space contains extensive high-temperature regions where thermal fluctuations above the effective Curie temperature T_C collapse the magnetization texture into an indistinguishable paramagnetic noise regime, regardless of the specific values of $\tilde{J}_2$ or $\tilde{K}_{DM}$. Within these thermally disordered regimes, configurations generated by

widely different parameter values relax into statistically similar noise textures, inflating the absolute prediction error even when the global correlation structure, captured by R^2 , remains strong. The MAE values for $T^{(0)}$, $\tilde{J}_2$, and $\tilde{K}_{DM}$ therefore reflect the thermal indistinguishability of the high-temperature paramagnetic regime rather than a systematic modeling deficiency: the same parameter that is most accurately recovered in the ordered low-temperature regime (where the spin texture preserves a clear topological signature) becomes the least constrained in the disordered high-temperature regime (where the texture collapses into thermal noise), producing non-negligible global MAE values driven primarily by the ill-conditioned, temperature-dominated subset of the dataset. A rigorous quantitative decomposition of this effect requires first separating the dataset into its constituent magnetic structural categories, as the global aggregation of all configurations inevitably conflates high-accuracy ordered-phase predictions with fundamentally ill-posed paramagnetic-regime estimates. This per-phase analysis is addressed further ahead, where the unsupervised structural clustering of the magnetization images allows us to isolate and quantify the specific contribution of each magnetic phase to the observed prediction error.

Table 3: Per-variable R^2 and MAE across the five evaluated architectures (Xception, DenseNet, ResNet, ViT, and the conventional CNN baseline) on the held-out test set. Each parameter shows R^2 (top row) and MAE (bottom row, in its physical units: K for $T^{(0)}$, meV for exchange/DMI, meV/atom for anisotropy/field); the best R^2 per parameter is in bold. Negative R^2 indicates predictions no better than the dataset mean.

Param.		Xcep.	Dense.	ResNet	ViT	CNN
$T^{(0)}$	R^2	0.96	0.95	0.94	0.95	0.94
	MAE	0.34	0.39	0.42	0.39	0.49
$\tilde{J}_2$	R^2	0.89	0.87	0.88	0.89	0.88
	MAE	0.018	0.021	0.018	0.018	0.019
$\tilde{J}_3$	R^2	0.01	-0.21	-0.24	-0.29	0.00
	MAE	0.030	0.031	0.032	0.032	0.028
$\tilde{J}_4$	R^2	0.00	-0.06	-0.67	-0.30	0.00
	MAE	0.022	0.024	0.028	0.025	0.022
$\tilde{K}_{an1}$	R^2	0.94	0.90	0.92	0.92	0.86
	MAE	0.052	0.057	0.054	0.055	0.060
$\tilde{K}_{anS}$	R^2	0.64	0.62	0.66	0.63	0.68
	MAE	0.027	0.028	0.027	0.027	0.026
$\tilde{H}_{ex}$	R^2	0.96	0.95	0.96	0.97	0.95
	MAE	0.013	0.013	0.012	0.012	0.013
$\tilde{K}_{DM}$	R^2	0.99	0.99	0.99	0.99	0.99
	MAE	0.019	0.032	0.032	0.037	0.028

Conversely, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ delimit the fundamental identifiability boundary of the inverse problem in the 2D observation space. They constitute the only genuinely unidentifiable parameters: their R^2 scores remain near zero across all evaluated architectures, indicating that the models are unable to recover meaningful information about them beyond the dataset mean. Physically, the spatial signatures of these subordinate exchange constants are energetically dominated by the primary $\tilde{J}_1$ and $\tilde{J}_2$ interactions, masking their individual effect on the magnetization texture and rendering them effectively unidentifiable from the s_z projection alone. Notably, their MAE values, on the order of 0.03 meV, remain numerically compa-

1424 rable to those of the better-identified parameters, which is a
 1425 direct consequence of their narrow physical range rather than
 1426 genuine predictive skill: a model defaulting to the dataset mean for $\tilde{J}_3$ and $\tilde{J}_4$ naturally incurs a small absolute error simply be-
 1427 cause the dynamic range of these parameters is itself limited. In
 1428 contrast, and in marked departure from what a purely in-plane
 1429 field geometry would permit, the external Zeeman field $\tilde{H}_{ex}$
 1430 and the bulk anisotropy $\tilde{K}_{an1}$ are recovered with high fidelity
 1431 ($R^2 \gtrsim 0.94$). Applying the field along the out-of-plane axis
 1432 in part of the dataset couples the Zeeman term directly to the
 1433 observable s_z component, while the extended anisotropy range
 1434 imprints sufficiently distinct domain-wall morphologies for $\tilde{K}_{an1}$
 1435 to leave a strong, recoverable spatial signature. The surface
 1436 anisotropy $\tilde{K}_{anS}$ occupies an intermediate regime ($R^2 \approx 0.65$),
 1437 as its visual manifestation remains partially entangled with the
 1438 bulk anisotropy $\tilde{K}_{an1}$ and is geometrically confined to the thin
 1439 boundary layer of the nanodot. **Taken together, these results
 1440 establish that six of the eight Hamiltonian parameters are effec-
 1441 tively recovered from the two-dimensional s_z projection, with
 1442 only the two higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ marking
 1443 the fundamental identifiability limit of the inverse problem.**

1444 To gain deeper insight into the internal representations learned
 1445 by the best-performing model, Xception, during the inverse
 1446 Hamiltonian estimation task, Figure ?? presents a Uniform Man-
 1447 ifold Approximation and Projection (UMAP) of the model's
 1448 latent space [?], where each point corresponds to a test sample
 1449 colored by the value of each of the eight Hamiltonian param-
 1450 eters θ . The UMAP embedding was computed from the high-
 1451 dimensional feature vectors extracted from the Xception back-
 1452 bone prior to the regression head, providing a two-dimensional
 1453 visualization of the geometric structure that the model has
 1454 learned to impose on the magnetization image manifold. Rather
 1455 than forming a single undifferentiated cloud, the latent space
 1456 exhibits a rich topological structure composed of several geo-
 1457 metrically distinct clusters and continuous gradients, reflecting
 1458 the fact that the model has organized the magnetization textures
 1459 according to their underlying physical regimes. Examining the
 1460 coloring by $T^{(0)}$ reveals the clearest and most spatially coherent
 1461 gradient across the embedding: low-temperature ordered config-
 1462 urations concentrate in well-separated compact regions, while
 1463 high-temperature disordered states occupy a large, diffuse clus-
 1464 ter in the upper portion of the latent space, consistent with the
 1465 strong predictive performance of the model on $T^{(0)}$. The DMI
 1466 strength $\tilde{K}_{DM}$ also produces a structured gradient in the latent
 1467 space, with high- $\tilde{K}_{DM}$ helical configurations forming a distinct
 1468 elongated region separated from the low- $\tilde{K}_{DM}$ disordered states
 1469 confirming that Xception has successfully encoded the chiral
 1470 winding imposed by the DMI as a geometrically meaningful
 1471 direction in its latent representation. The bulk anisotropy $\tilde{K}_{an1}$
 1472 and the external field $\tilde{H}_{ex}$ likewise display coherent, structured
 1473 gradients across the embedding, mirroring their high regression
 1474 accuracy and confirming that Xception encodes them as well-
 1475 defined directions in its latent geometry, a direct consequence of
 1476 the out-of-plane field configurations and the extended anisotropy
 1477 range, which render these parameters observable in the s_z pro-
 1478 jection. In contrast, the longer-range exchange constants $\tilde{J}_3$ and
 1479 $\tilde{J}_4$ produce nearly uniform colorings across the entire embed-




Figure 6: UMAP-based 2D projection of the Xception latent space on the held-out test set, colored by each of the eight Hamiltonian parameters $\theta = [T^{(0)}, \tilde{J}_2, \tilde{J}_3, \tilde{J}_4, \tilde{K}_{an1}, \tilde{K}_{anS}, \tilde{H}_{ex}, \tilde{K}_{DM}]$. Each point represents a magnetization image $\mathbf{X}_n$ projected from the high-dimensional feature space of the Xception backbone into two dimensions via UMAP.

ding, indicating that the model's internal representations carry little discriminative information about these parameters, a result that is fully consistent with their low R^2 values and their energetic subordination to the dominant first-shell exchange $\tilde{J}_1 = 1.0$ meV. The surface anisotropy $\tilde{K}_{anS}$ occupies an intermediate case, showing a partially mixed coloring consistent with its moderate identifiability and its entanglement with the bulk anisotropy signature.

Beyond the per-parameter gradients, the same embedding reveals that Xception has organized the magnetization textures into geometrically distinct and physically interpretable regions, without any explicit supervision of magnetic phase labels. The systematic variation of $T^{(0)}$ separates ordered, low-temperature configurations, which concentrate in compact, well-separated regions, from the thermally disordered states that occupy a large, diffuse cluster at high $T^{(0)}$, where thermal fluctuations dominate over the competing energetic contributions of $\tilde{J}_1$ and $\tilde{K}_{DM}$. Superimposed on this temperature axis, the $\tilde{K}_{DM}$ gradient further resolves the ordered region into the chiral helical, dense skyrmionic, and collinear Ising-like regimes, while the mixed-chirality configurations form their own transitional zone. This spontaneous geometric separation motivates a structured clustering analysis, performed directly on the two-dimensional UMAP embedding of the test set feature vectors (min-max normalized to $[0, 1]$ along each axis), which we use next to identify and characterize the dominant magnetic phases.

To identify dominant magnetic phases encoded in the latent space, a density-based clustering analysis was performed using the HDBSCAN algorithm on the UMAP embedding, followed by a KMeans subdivision of anomalously large clusters. Based on visual inspection of representative samples and the mean values of the Hamiltonian parameters θ within each group, the resulting clusters were consolidated into five physically meaningful magnetic phase categories.

The *helical phase* contains the most spatially coherent configurations: well-defined stripe and labyrinthine patterns with consistent chiral winding and periodicity, stabilized at low temperature by a moderate DMI ($\langle \tilde{K}_{DM} \rangle \approx 0.75$ meV). The *skyrmion phase* groups configurations in which a strong DMI ($\langle \tilde{K}_{DM} \rangle \approx 0.88$ meV) stabilizes dense lattices of isolated, radially symmetric solitons, visible as arrays of reversed- s_z cores in a uniform background. The *Ising phase* comprises low-temperature ($\langle T^{(0)} \rangle \approx 1.86$ K), high-anisotropy configurations with sharp,

discrete domain walls between collinear up/down regions, reflecting the extended anisotropy range. The *bimeron phase* gathers mixed-chirality textures in which in-plane and out-of-plane components coexist as paired meron-antimeron structures forming irregular domains without long-range periodicity. Finally, the *others* category is the most heterogeneous, gathering thermally disordered, transitional, and topologically mixed states whose common trait is the absence of a distinctive, reproducible pattern, making them the most ill-conditioned samples.

Figure ?? presents five representative magnetization images per phase, selected as the samples closest to the centroid of each group in the UMAP embedding, along with the mean Hamiltonian parameters $\langle T^{(0)} \rangle$, $\langle \tilde{J}_2 \rangle$, and $\langle \tilde{K}_{DM} \rangle$.




Figure 7: Representative magnetization image $\mathbf{X}_n$ for each of the five magnetic phase categories identified by clustering the Xception latent space (the sample nearest the cluster centroid in the UMAP embedding); color encodes the out-of-plane component m_z (+1 spin up, -1 spin down). The phases span ordered helical stripes, thermally disordered *others*, mixed-chirality bimerons, dense skyrmion lattices, and sharp Ising domain walls.

The per-phase regression performance of the best-performing model, Xception, is summarized in Table ?. This per-phase evaluation, reported separately for each of the five magnetic phases isolated by the latent-space clustering, probes how the model behaves across qualitatively distinct physical regimes rather than as a single homogeneous distribution. The results reveal a pronounced heterogeneity in predictive accuracy across magnetic phases that is physically consistent with the identifiability analysis presented earlier in this section.

The *skyrmion* and *Ising* phases achieve the highest R^2 values (typically 0.94–0.99) and the lowest MAE across essentially all identifiable parameters, reflecting the fact that these highly ordered configurations, dense soliton lattices and sharp, discrete domain walls, respectively, exhibit the most discriminative and spatially coherent signatures in the s_z projection. Within these phases, the well-defined periodicity and topological structure of the magnetization texture provide unambiguous observational constraints that allow the model to recover the underlying Hamiltonian parameters with high fidelity. The *helical* phase follows closely, with suitable recovery of $T^{(0)}$ and $\tilde{K}_{DM}$ ($R^2 \approx 0.99$) and good performance for $\tilde{J}_2$ and $\tilde{H}_{ex}$, although its anisotropy constants are estimated less accurately owing to the weaker anisotropic imprint in the chiral stripe regime.

In stark contrast, the *others* and *bimeron* categories concentrate the largest absolute errors across virtually all Hamiltonian parameters, as evidenced by the dominant MAE values in Table ?. This result directly confirms the physical interpretation advanced earlier: the thermally disordered configurations grouped in the *others* category, and the mixed-chirality meron-antimeron textures of the *bimeron* phase, relax into structurally weakly-discriminative magnetization patterns, so that widely different parameter combinations produce statistically similar textures

Table 4: Per-phase R^2 and MAE of Xception across the Hamiltonian parameters θ for the five magnetic phases. Each parameter shows R^2 (top row) and MAE (bottom row, in its physical units); the best R^2 per parameter is in bold. N/A denotes cases with near-zero target variance where R^2 is undefined, and a dash (–) indicates that no valid metric is available.

Param.		Helical	Others	Bimer.	Skyrm.	Ising
$T^{(0)}$	R^2	0.99	0.88	0.83	0.99	0.99
	MAE	0.18	0.84	0.99	0.18	0.14
$\tilde{J}_2$	R^2	0.76	0.66	0.62	0.96	N/A
	MAE	0.001	0.089	0.028	0.001	–
$\tilde{J}_3$	R^2	N/A	0.01	0.03	N/A	N/A
	MAE	–	0.15	0.13	–	–
$\tilde{J}_4$	R^2	N/A	–0.00	0.00	N/A	N/A
	MAE	–	0.12	0.10	–	–
$\tilde{K}_{an1}$	R^2	0.47	0.92	0.15	0.94	0.98
	MAE	0.047	0.018	0.029	0.10	0.14
$\tilde{K}_{anS}$	R^2	0.54	0.43	0.34	N/A	N/A
	MAE	0.030	0.019	0.025	–	–
$\tilde{H}_{ex}$	R^2	0.83	N/A	0.54	0.97	0.93
	MAE	0.009	–	0.002	0.021	0.005
$\tilde{K}_{DM}$	R^2	0.99	0.99	0.73	0.97	0.96
	MAE	0.020	0.004	0.014	0.026	0.022

tures and inflate the inverse-mapping error within these regimes.

The elevated MAE for $T^{(0)}$ in the *others* phase is particularly revealing, as it quantitatively isolates the contribution of the high-temperature paramagnetic regime to the global prediction error observed in the aggregated analysis; likewise, the *bimeron* phase yields the lowest R^2 for the anisotropy and field parameters, reflecting the intrinsic spatial ambiguity of its coexisting in-plane and out-of-plane components. Across all five phases, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ retain near-zero R^2 , confirming at the per-phase level that their unidentifiability is a fundamental property of the s_z observation rather than an artifact of any particular structural regime.

5.3. RAMs-based Interpretability Results

To spatially trace the predictive behavior of the Xception architecture across the identified magnetic phase categories, RAMs were computed following the CNN-based formulation introduced in Section ?. For each of the five magnetic phases a single representative magnetization image $\mathbf{X}_n$ was selected as the sample closest to the centroid of the corresponding cluster group in the UMAP latent space, maximizing the concentration of gradient activations within the physical extent of the nanodot disk. For the spatial overlays, RAMs were computed for three representative Hamiltonian parameters ($T^{(0)}$, $\tilde{J}_2$, and $\tilde{K}_{DM}$) across eight strategically selected convolutional layers of the Xception backbone, spanning the full depth of the network: block1_conv2: entry-level features at spatial resolution 112×112 , capturing low-level pixel-scale magnetization gradients. block3_sepconv2 and block4_sepconv2: intermediate entry-flow layers at 56×56 and 28×28 , encoding domain boundary frequencies and local spin-winding patterns. block6_sepconv2 through block12_sepconv2: middle-flow layers at 14×14 , constituting the deepest repeated processing stage of Xception and encoding progressively more

abstract topological representations of the magnetic texture
block13_sepconv2: exit-flow transition layer prior to global
pooling, encoding high-level semantic representations of the
global magnetic phase.

Now, the precision factor γ in the RAM objective $O_p = \exp(-\gamma(\theta_p - \hat{\theta}_p)^2)$ governs the selectivity of the gradient signal: small values render the objective nearly constant regardless of prediction quality, producing spatially uninformative maps, while excessively large values collapse the gradient to zero for parameters with non-negligible prediction error, yielding empty activation maps. To select γ in a principled and data-driven manner, a sensitivity analysis was conducted by evaluating the objective O_p at the empirical mean prediction error $\bar{\varepsilon}_p = \langle |\theta_p - \hat{\theta}_p| \rangle$ of each Hamiltonian parameter on the test set, for values spanning $\gamma \in \{0.1, 0.3, 0.5, 1, 2, 5, 10, 20, 50, 100\}$.

Figure ?? illustrates the resulting sensitivity curves. The left panel shows the shape of the Gaussian objective as a function of the prediction error ε for each candidate γ : values below $\gamma = 1$ produce flat curves that activate almost uniformly regardless of prediction quality, while values above $\gamma = 20$ generate overly steep curves that suppress the gradient for errors as small as $\varepsilon \approx 0.2$; the selected $\gamma = 10$ provides an intermediate, well-graded response. The right panel evaluates O_p at the empirical mean error of each parameter as a function of γ : because the identifiable parameters ($T^{(0)}$, $\tilde{J}_2$, and $\tilde{K}_{DM}$) are recovered with small mean errors, their objective remains close to unity across the entire γ range and never approaches the gradient collapse threshold $O_p < 0.05$. Only the poorly identified higher-order exchange constant $\tilde{J}_3$, whose mean error is comparatively large, shows a noticeable decline at high γ (down to $O_p \approx 0.8$ at $\gamma = 100$), yet even it stays well above the collapse regime. The value $\gamma = 10$ was therefore selected as the operating point: it is selective enough to yield spatially discriminative maps (steep left-panel response) while keeping a well-conditioned gradient signal for all Hamiltonian parameters.

To ensure that activation magnitudes remain physically interpretable and spatially confined to the nanodot geometry, each RAM was masked by a circular binary mask aligned with the disk boundary prior to normalization. The globally normalized RAMs $\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)}$ were obtained following Equation ??, and the spatial patterns were additionally visualized under local per-cell normalization to reveal the gradient attribution structure at each individual layer.

Figure ?? presents the resulting RAM overlays for the helical phase, selected for visualization because its coherent stripe and labyrinthine patterns, with consistent chiral winding and periodicity, produce the clearest and most physically interpretable gradient attribution across network depth. Three representative layers are shown, spanning the full depth of the backbone: the entry-flow layer block1_conv2, a middle-flow layer block6_sepconv2, and the terminal exit-flow layer block13_sepconv2. The spatial attribution patterns reveal a physically consistent progression. In the early layer, the activations for $\tilde{K}_{DM}$ are spatially dense and sharply structured, tracing the chiral domain walls and stripe boundaries that encode the DMI strength. As depth increases, these activations progres-

sively consolidate into coarser, spatially localized hotspots along the ordered stripe pattern, while the $T^{(0)}$ attribution spreads over the broader texture that reflects the global degree of thermal order. In the terminal layer, activations become concentrated in a small number of compact regions, indicating that the network's final semantic representations encode global topological stability rather than local texture periodicity. Consistently, the $\tilde{J}_2$ row remains nearly inactive across all layers, in agreement with its vanishing value in this phase. This hierarchical progression matches the theoretical prediction, namely, early layers resolve the short-range energetic competition between $\tilde{J}_1$ and $\tilde{K}_{DM}$, while deeper layers integrate these local features into macro-scale thermodynamic representations governed by $T^{(0)}$.

To quantify how each network layer contributes to the estimation of each parameter, Figure ?? presents the mean maximum RAM activation per layer and variable, averaged over $N = 300$ randomly sampled test images per magnetic phase and displayed as a heatmap. Because the absolute activation magnitude differs across parameters, the values are normalized *row-wise*: for each parameter, the eight layer values are standardized via a z-score and mapped through a sigmoid, $\sigma\left(\frac{x-\mu}{\sigma_x}\right) \in (0, 1)$, where μ and σ_x are the mean and standard deviation of that parameter's activations across layers. This normalization reveals, for each parameter independently, the relative layer depth at which its gradient response is strongest (warm cells) versus weakest (cool cells), rather than comparing magnitudes between parameters.

This analysis confirms and extends the qualitative observations from the spatial maps. For the identifiable parameters ($T^{(0)}$, $\tilde{J}_2$, $\tilde{K}_{DM}$, $\tilde{K}_{an1}$, and $\tilde{H}_{ex}$), the layer profile is markedly non-monotonic, with a consistently dominant warm cell at the terminal exit-flow layer block13_sepconv2, where the network aggregates the high-level semantic representations of these parameters. Secondary warm cells in the middle-flow layers (block6–block8) reflect the intermediate stage at which the localized chiral and domain-wall features are assembled, and are most pronounced in the ordered skyrmion and Ising phases whose sharp textures provide the strongest spatial signal. By contrast, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ display a flat, unstructured layer profile with no layer standing out, indicating that no network depth resolves a distinctive spatial feature for them, consistent with their fundamental unidentifiability established in Section ?? and with their energetic subordination to the dominant first- and second-shell exchange. The identifiable anisotropy and field parameters, in turn, show their clearest depth structure in the phases where they leave a measurable imprint (the high-anisotropy Ising regime and the out-of-plane-field configurations), consistent with their successful recovery in the regression analysis.

The RAM analysis thus isolates a physically coherent set of features as the primary drivers of accurate regression: chiral domain-wall boundaries, stripe periodicity, and skyrmion cores. Conversely, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ never acquire spatially concentrated activations, confirming that their limited identifiability is not a consequence of insufficient model capacity but of the fundamental physical constraints of the inverse problem: their energetic subordination to the dominant





Figure 8: Sensitivity analysis of the RAM precision factor γ . Left: shape of the Gaussian objective $O_p = e^{-\gamma \epsilon^2}$ versus the prediction error $\epsilon = |\theta_p - \hat{\theta}_p|$ for $\gamma \in \{0.1, \dots, 100\}$ (selected $\gamma = 10$ in red; vertical dotted line = empirical mean error). Right: O_p evaluated at each parameter’s mean error as a function of γ ; the horizontal dotted line marks the gradient-collapse threshold $O_p = 0.05$. The value $\gamma = 10$ is the selected operating point.





Figure 10: Heatmap of the mean maximum RAM activation per network layer (columns, B1–B13) and Hamiltonian parameter θ (rows), for the five magnetic phases, averaged over $N = 300$ test samples per phase ($N = 211$ for Ising; Xception CNN) and measured within the nanodot disk (masked RAM). Activations are normalized row-wise (per parameter, across layers) via a sigmoid $\sigma(\frac{z-\mu}{\sigma_x})$, so warm cells mark the layers where each parameter’s gradient response is relatively strongest. The identifiable parameters peak at the exit-flow layer B13, whereas $\tilde{J}_3$ and $\tilde{J}_4$ show a flat, unstructured profile with no dominant layer.

Figure 9: RAMs $\tilde{\mathcal{R}}_{\text{RAM}}^{(l,p)}$ for the Xception architecture in the helical phase. Rows are three Hamiltonian parameters $\{T^{(0)}, \tilde{J}_2, \tilde{K}_{DM}\}$; columns are three representative network layers, block1 (entry, B1), block6 (middle, B6), and block13 (exit, B13), with the original s_z image in the leftmost column. Each panel overlays the RAM intensity (locally normalized) on the magnetization; the value below each original image is the ground-truth parameter.

first- and second-shell exchange, which leaves no distinguishable imprint on the s_z projection. For completeness, the RAM formulation was also derived for the Vision Transformer (Section ??) since Xception is the best-performing model and provides the clearest spatial attributions, the remainder of the interpretability analysis focuses on it.

5.4. Interpretability Benchmark: RAMs vs. Baselines

Having established that RAMs isolate physically meaningful spatial features, we now benchmark them quantitatively against

the perturbation-based attribution baselines introduced in Section ??, LIME, SHAP, and SF-SHAP, using the deletion protocol of Equations ??–??. The comparison targets both fidelity (how faithfully each map localizes the pixels that drive the prediction) and computational cost.

The perturbation baselines operate on superpixels rather than individual pixels, so the spatial granularity of the segmentation directly conditions their behavior. To fix this granularity on a physical basis, the number of SLIC segments was chosen so that the resulting superpixel equivalent diameter matches the characteristic magnetic domain width of the dataset (≈ 4.9 px). As shown in Figure ??, this criterion is satisfied at $N = 100$ segments, which we adopt for all perturbation baselines, ensuring that each interpretable region corresponds to one coherent magnetic domain rather than to an arbitrary pixel cluster.

Figure ?? reports the deletion curves for the three represen-




Figure 11: Physical criterion for the SLIC segmentation used by the perturbation-based baselines. The superpixel equivalent diameter is plotted against the number of SLIC segments; the selected value $N = 100$ (dashed red) yields superpixels whose size matches the characteristic magnetic domain width (≈ 4.9 px, dashed green), so that each interpretable region corresponds to a single magnetic domain.

tative parameters. Two consistent trends emerge. First, for spatially localized parameters, most notably $\tilde{K}_{DM}$, all methods, including RAMs, degrade the prediction far faster than the random-relevance baseline; RAMs thus genuinely localize the pixels carrying the DMI signal, at a deletion AUC of the same order as the perturbation methods. Second, for the global parameter $T^{(0)}$ the situation reverses: no method separates from the random baseline. This is physically expected rather than a shortcoming: the annealing temperature is not tied to any specific domain sector but is an *ambient* property of the overall degree of order, so distributed (random) deletions perturb it more effectively than localized ones. The deletion metric thus correctly reflects that $T^{(0)}$ is a delocalized, environment-like parameter, in contrast to the spatially bound $\tilde{J}_2$ and $\tilde{K}_{DM}$.

The advantage of RAMs lies in their computational cost. Figure ?? contrasts the cost of producing a single relevance map. Being gradient-based, a RAM requires exactly one backward pass (i.e., a single model evaluation) whereas the perturbation baselines must query the model over many masked coalitions: of order 10^3 evaluations for LIME and SHAP and $2^7 = 128$ for SF-SHAP. This translates into a measured wall-clock cost per map that is one to two orders of magnitude lower for RAMs than for any perturbation method. RAMs therefore deliver a faithfulness competitive with established attribution techniques while remaining tractable for the high-throughput, per-sample interpretability that materials-characterization pipelines require.

5.5. Limitations

The most fundamental limitation of the framework is the partial observability of the extended Heisenberg Hamiltonian from two-dimensional magnetization images. Indeed, projecting the three-dimensional spin configuration onto the out-of-plane component s_z of the central nanodot layer irreversibly discards physical information. While the dataset enrichment adopted in this work, the out-of-plane field configurations and the extended anisotropy range, renders six of the eight parameters recoverable, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ remain fundamentally unidentifiable: their energetic contribution is subordinated

to the dominant first- and second-shell exchange, so they leave no distinguishable signature in the s_z texture, regardless of model capacity or architecture. A related, milder limitation affects the surface anisotropy $\tilde{K}_{anS}$, which is recovered only partially because the input images are extracted exclusively from the central discrete layer of the three-dimensional nanodot lattice, discarding the depth-resolved spin structure that encodes the surface-to-bulk anisotropy ratio $\tilde{K}_{anS}/\tilde{K}_{an1}$. Recovering these remaining parameters with high fidelity would require either access to the full three-dimensional spin vector field at each lattice site, or volumetric magnetization inputs obtained through multi-layer or tomographic magnetic imaging techniques. This is not a limitation of the deep learning models themselves but of the observational modality assumed by the framework, which mimics the two-dimensional projection inherent to most experimental magnetic imaging techniques such as magnetic force microscopy or Lorentz transmission electron microscopy.

Besides, the dataset consists entirely of simulated magnetization images. Although the simulation incorporates realistic finite-size effects, cylindrical boundary conditions, and thermodynamic relaxation, it does not model several experimental phenomena, thermal noise and instrument point-spread functions, crystallographic defects and grain boundaries, surface reconstruction, and substrate-induced strain, that can renormalize the effective Hamiltonian parameters in real nanodot imaging. Consequently, a model trained only on simulated data may degrade on experimentally acquired images, a domain-shift problem not quantified here; bridging this simulation-to-experiment gap would require transfer learning from a small set of labeled experimental samples or physics-informed data augmentation replicating experimental noise.

A further limitation concerns the internal structure of the sampled parameter space. The dataset spans two mechanisms for stabilizing modulated textures, exchange frustration (driven by $\tilde{J}_2$) and the chiral DMI (driven by $\tilde{K}_{DM}$), which were deliberately activated in a mutually exclusive fashion to isolate and compare their individual contributions, and were subsequently unified into a single training corpus. While this design cleanly separates the two mechanisms and demonstrates that both can independently generate analogous modulated states, it does not include configurations in which the two interactions coexist and compete simultaneously, as may occur in real materials. Extending the dataset to jointly sample $\tilde{J}_2$ and $\tilde{K}_{DM}$ would allow the model to learn their combined effect.

Finally, RAMs share the limitations of all gradient-based saliency methods: the maps reflect the local sensitivity of the objective O_p at the point of evaluation rather than a globally faithful explanation, and gradient saturation in deep networks can yield diffuse attributions for parameters with low predictive signal, precisely the unidentifiable ones. Moreover, the current formulation produces a single map per image per parameter without capturing attribution uncertainty; ensemble-based or Bayesian extensions could supply confidence intervals.




Figure 12: Deletion curves (Most Relevant First) for RAM, LIME, SHAP, and SF-SHAP, with a random-relevance lower bound, for the three representative parameters $T^{(0)}$, J_2 , and K_{DM} . The vertical axis is the induced deviation $|\hat{\theta}_{pert} - \hat{\theta}_0|$ as the most relevant pixels are progressively replaced by the baseline; the area under each curve (AUC, in the legend) quantifies faithfulness.




Figure 13: Computational cost of the interpretability methods. Left: measured wall-clock time per relevance map (log scale). Right: number of model evaluations per map (log scale). A RAM requires a single backward pass, against $\sim 10^3$ evaluations for LIME and SHAP and 128 for SF-SHAP.

6. Conclusions

This work introduced a regression-based explainable deep learning framework for the inverse estimation of Hamiltonian parameters from two-dimensional magnetization images of magnetic nanodots, integrating GPU-accelerated atomistic Monte Carlo simulations of an eight-parameter extended Heisenberg Hamiltonian, five deep learning architectures, an unsupervised UMAP/HDBSCAN latent-space analysis, and RAMs, a novel spatial interpretability formulation that generalizes gradient-based activation mapping to continuous regression.

Among the five evaluated architectures, Xception was the most accurate, although all models, including the compact CNN baseline trained from scratch, performed comparably, showing that accuracy is governed by the physical information in the S_z projection rather than by model capacity or pretraining. Six of the eight parameters were recovered with high fidelity ($T^{(0)}$, J_2 , K_{DM} , $\tilde{K}_{an1}$, $\tilde{H}_{ex}$, with $\tilde{K}_{anS}$ at an intermediate level), enabled by the dataset enrichment with out-of-plane field configurations and an extended anisotropy range. By contrast, the higher-order exchange constants $\tilde{J}_3$ and $\tilde{J}_4$ retained near-zero R^2 across all architectures, establishing that the limitation is not model capacity but the intrinsic partial observability of these subordinate interactions from the two-dimensional projection.

Beyond the aggregate regression performance, the UMAP projection of the Xception latent representations revealed that the model autonomously organized the magnetization textures into

geometrically distinct and physically interpretable regions of the embedding space, without any explicit supervision of magnetic phase labels. Density-based clustering identified five dominant magnetic phase categories: helical, skyrmion, Ising, bimeron, and a heterogeneous *others* group. The per-phase regression analysis confirmed that the model achieves its highest accuracy within the most ordered phases (skyrmion and Ising, $R^2 \gtrsim 0.94$ across the identifiable parameters), where the magnetization texture carries the most discriminative physical information, while the disordered and mixed-chirality configurations concentrate the largest errors.

The proposed RAMs formulation extended gradient-based activation mapping to continuous regression through a scale-invariant Gaussian error-decay objective $O_p = \exp(-\gamma_p(\theta_p - \hat{\theta}_p)^2)$, yielding physically interpretable attribution maps per parameter. The layer-resolved analysis revealed a hierarchical decoupling across network depth: early layers capturing fine-grained domain periodicity and chiral winding, deeper layers consolidating them into global representations, with the DMI attributions localizing on the features that encode it. Benchmarked against LIME, SHAP, and SF-SHAP through a deletion protocol, RAMs attained competitive spatial faithfulness at a single backward pass per map, one to two orders of magnitude cheaper than the perturbation-based methods, making them practical for per-sample interpretability in high-throughput characterization.

Taken together, these results demonstrate that regression-based deep learning, combined with physically motivated interpretability, constitutes a viable and physically meaningful approach to inverse Hamiltonian estimation in magnetic nanostructures. The RAMs framework is architecture-agnostic and directly applicable to any continuous regression task requiring spatial attribution of physical parameters, extending beyond magnetic nanodots to other spatially resolved imaging modalities.

Several directions merit future investigation: (i) using the full three-dimensional spin field or volumetric inputs to relax the partial-observability constraint and recover the elusive $\tilde{K}_{anS}$, $\tilde{J}_3$, and $\tilde{J}_4$; (ii) domain adaptation and few-shot transfer learning to bridge the simulation-to-experiment gap; (iii) uncertainty quantification of the RAM maps via ensemble or Bayesian attribution; and (iv) active-learning strategies that select the most informative configurations to simulate, accelerating dataset construction

for higher-dimensional Hamiltonian spaces.

Acknowledgments

Under grants provided by the research project: "Aprendizaje automático aplicado a la predicción eficiente de propiedades magnéticas de nanopartículas como soporte a la generación de dispositivos electrónicos" Hermes 62642, funded by Universidad Nacional de Colombia. Also, A. Alvarez thanks the project: "Estimación de la capacidad de generación de energía eléctrica del agua coproducida en campos de petróleo y gas en Colombia a partir de técnicas de aprendizaje automático informado por la física," Code 111908, funded by 951-2024 Convocatoria fortalecimiento del conocimiento científico y tecnológico de las fuentes - Minciencias.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author(s) used Gemini and Claude in order to enhance the writing style. After using this tool/service, the author(s) reviewed and edited the content as needed and take(s) full responsibility for the content of the published article.

References

- [1] A. Heinrich, W. Oliver, L. Vandersypen, A. Ardavan, R. Sessoli, D. Loss, A. Jayich, J. Fernandez, A. Laucht, and A. Morello, "Quantum-coherent nanoscience," *Nat. Nanotechnol.*, vol. 16, no. 12, pp. 1318–1329, 2021.
- [2] A. Fert, R. Ramesh, V. Garcia, F. Casanova, and M. Bibes, "Electrical control of magnetism by electric field and current-induced torques," *Rev. Mod. Phys.*, vol. 96, no. 1, p. 015005, 2024.
- [3] X. Battle, C. Moya, M. Escoda, O. Iglesias, A. Fraile, and A. Labarta, "Magnetic nanoparticles: from the nanostructure to the physical properties," *J. Magn. Magn. Mater.*, vol. 543, p. 168594, 2022.
- [4] I. Charalampidis and J. Barker, "Metadynamics calculations of the effect of thermal spin fluctuations on skyrmion stability," *arXiv preprint arXiv:2310.03169*, 2023.
- [5] X. Yu, "Magnetic imaging of various topological spin textures and their dynamics," *J. Magn. Magn. Mater.*, vol. 539, p. 168332, 2021.
- [6] Y. Tokura and N. Kanazawa, "Magnetic skyrmion materials," *Chem. Rev.*, vol. 121, no. 5, pp. 2857–2897, 2020.
- [7] V. Kuchkin, B. Barton, F. Rybakov, S. Blügel, B. Schroers, and N. Kiselev, "Magnetic skyrmions, chiral kinks, and holomorphic functions," *Phys. Rev. Lett.*, vol. 125, no. 18, p. 187202, 2020.
- [8] J. Kong, Y. Ren, M. Tey, P. Ho, K. Khoo, X. Chen, and A. Soumyanarayanan, "Quantifying the magnetic interactions governing chiral spin textures using deep neural networks," *ACS Appl. Mater. Interfaces*, vol. 16, no. 1, pp. 1025–1032, 2023.
- [9] M. Goerzen, S. Malottki, G. Kwiatkowski, P. Bessarab, and S. Heinze, "Atomistic spin simulations of electric-field-assisted nucleation and annihilation of magnetic skyrmions in Pd/Fe/Ir(111)," *Phys. Rev. B*, vol. 105, no. 21, p. 214435, 2022.
- [10] A. Aldarawsheh, M. Sallermann, M. Abusaa, and S. Lounis, "Intrinsic néel antiferromagnetic multimeronic spin textures in ultrathin films," *J. Phys. Chem. Lett.*, vol. 14, no. 40, pp. 8970–8978, 2023.
- [11] C. Back, V. Cros, H. Ebert, K. Everschor, A. Fert, M. Garst, T. Ma, S. Mankovsky, T. Monchesky, and M. Mostovoy, *et al.*, "The 2020 skyrmionics roadmap," *J. Phys. D: Appl. Phys.*, vol. 53, no. 36, p. 363001, 2020.
- [12] R. Murakami, M. Mizumaki, I. Akai, and H. Shouno, "Inverse estimation of parameters for the magnetic domain via dynamics matching using visual-perceptive similarity," *Sci. Technol. Adv. Mater.: Methods*, vol. 2, no. 1, pp. 139–152, 2022.
- [13] H. Kwon, H. Yoon, C. Lee, G. Chen, K. Liu, A. Schmid, Y. Wu, J. Choi, and C. Won, "Magnetic hamiltonian parameter estimation using deep learning techniques," *Sci. Adv.*, vol. 6, no. 39, p. eabb0872, 2020.
- [14] T. Nguyen, L. Rasabathina, O. Hellwig, A. Sharma, G. Salvan, S. Yochelis, Y. Paltiel, L. Baczewski, and C. Tegenkamp, "Cooperative effect of electron spin polarization in chiral molecules studied with non-spin-polarized scanning tunneling microscopy," *ACS Appl. Mater. Interfaces*, vol. 14, no. 33, pp. 38013–38020, 2022.
- [15] J. Zhao, L. Bai, S. Li, Z. Cao, Y. Peng, J. Bai, X. Cai, X. Shi, X. Lin, and G. Wei, *et al.*, "An overview of advanced instruments for magnetic characterization and measurements," *Frontiers in Electronics*, vol. 6, p. 1645594, 2025.
- [16] J. Gaida, H. Lourenço, S. Yalunin, A. Feist, M. Sivilis, T. Hohage, F. García de Abajo, and C. Ropers, "Lorentz microscopy of optical fields," *Nat. Commun.*, vol. 14, no. 1, p. 6545, 2023.
- [17] M. d'Aquino, S. Perna, M. Pancaldi, R. Hertel, S. Bonetti, and C. Serpico, "Micromagnetic study of inertial spin waves in ferromagnetic nanodots," *Phys. Rev. B*, vol. 107, no. 14, p. 144412, 2023.
- [18] S. Lepadatu, "Micromagnetic monte carlo method with variable magnetization length based on the landau–lifshitz–bloch equation for computation of large-scale thermodynamic equilibrium states," *J. Appl. Phys.*, vol. 130, no. 16, 2021.

- [19] X. Yin, *Computational Materials Design: Integrating Physics With Optimization and Machine Learning*. PhD thesis, Carnegie Mellon University, 2023.
- [20] S. Karnatov, "Analysis of psnr, ssim, lpips metrics in the context of human perception of visual similarity," *Transport systems and technologies*, no. 46, 2025.
- [21] W. Lee, T. Song, and K. Kim, "Deep learning methods for hamiltonian parameter estimation and magnetic domain image generation in twisted van der waals magnets," *Mach. Learn.: Sci. Technol.*, vol. 5, no. 2, p. 025073, 2024.
- [22] Y. Chen, Q. Yang, Y. Li, H. Zhang, and C. Zhang, "Data-driven deep convolutional neural networks for electromagnetic field estimation of transformers," *IEEE Trans. Appl. Supercond.*, vol. 34, no. 8, pp. 1–5, 2024.
- [23] S. Lan, S. Li, and B. Shahbaba, "Scaling up bayesian uncertainty quantification for inverse problems using deep neural networks," *SIAM/ASA J. Uncertain. Quantif.*, vol. 10, no. 4, pp. 1684–1713, 2022.
- [24] L. Huang, J. Li, X. Ding, Y. Zhang, H. Chen, and A. Ozcan, "Cycle-consistency-based uncertainty quantification of neural networks in inverse imaging problems," *Intell. Comput.*, vol. 2, p. 0071, 2023.
- [25] M. Baldan, P. Di Barba, and D. Lowther, "Physics-informed neural networks for inverse electromagnetic problems," *IEEE Trans. Magn.*, vol. 59, no. 5, pp. 1–5, 2023.
- [26] W. Wang, Z. Wang, Y. Zhang, B. Sun, and K. Xia, "Learning order parameters from videos of skyrmion dynamical phases with neural networks," *Phys. Rev. Appl.*, vol. 16, no. 1, p. 014005, 2021.
- [27] T. Čadež and K. Kim, "Neural scaling laws for deep regression on domain image data of twisted magnets," *Mach. Learn.: Sci. Technol.*, vol. 7, no. 2, p. 025011, 2026.
- [28] M. He, B. Li, and S. Sun, "A survey of class activation mapping for the interpretability of convolution neural networks," in *Proc. Int. Conf. Signal Inf. Process., Netw. Comput.*, pp. 399–407, Springer, 2022.
- [29] J. Zeng, M. Albooyeh, M. Rajaei, A. Sifat, E. Potma, H. Wickramasinghe, and F. Capolino, "Direct detection of photoinduced magnetic force at the nanoscale reveals magnetic nearfield of structured light," *Sci. Adv.*, vol. 8, no. 45, p. eadd0233, 2022.
- [30] G. Silva, A. Bergeron, J. Plascak, and D. Landau, "Three-dimensional heisenberg model with dzyaloshinskii-moriya interaction: A monte carlo study," *Phys. Rev. E*, vol. 109, no. 6, p. 064113, 2024.
- [31] A. Talapatra, U. Gajera, S. P. J. Chelvane, and J. Mohanty, "Understanding the magnetic microstructure through experiments and machine learning algorithms," *ACS Appl. Mater. Interfaces*, vol. 14, no. 44, pp. 50318–50330, 2022.
- [32] Z. Zhao, J. Ye, and Y. Bresler, "Generative models for inverse imaging problems: From mathematical foundations to physics-driven applications," *IEEE Signal Process. Mag.*, vol. 40, no. 1, pp. 148–163, 2023.
- [33] J. Agudelo, E. Restrepo, and D. Landau, "A study of magnetic domains in thin films from fm and afm interactions," *Physica A*, vol. 517, pp. 542–550, 2019.
- [34] V. Borisov, Y. Kvashnin, N. Ntallis, D. Thonig, P. Thunström, M. Pereiro, A. Bergman, E. Sjöqvist, A. Delin, and L. Nordström, *et al.*, "Heisenberg and anisotropic exchange interactions in magnetic materials with correlated electronic structure and significant spin-orbit coupling," *Phys. Rev. B*, vol. 103, no. 17, p. 174422, 2021.
- [35] O. Eriksson, A. Bergman, L. Bergqvist, and J. Hellsvik, *Atomistic spin dynamics: foundations and applications*. Oxford university press, 2017.
- [36] B. Cullity and C. Graham, *Introduction to magnetic materials*. John Wiley & Sons, 2011.
- [37] A. Bogdanov and C. Panagopoulos, "Physical foundations and basic properties of magnetic skyrmions," *Nat. Rev. Phys.*, vol. 2, pp. 492–498, 2020.
- [38] D. Landau and K. Binder, *A guide to Monte Carlo simulations in statistical physics*. Cambridge university press, 2021.
- [39] L. Desplat, *Thermal Stability of Metastable Magnetic Skyrmions*. Springer Theses, Springer International Publishing, 2021.
- [40] K. Murphy, *Probabilistic machine learning: Advanced topics*. MIT press, 2023.
- [41] B. Koonce, "Resnet 50," in *Convolutional neural networks with swift for tensorflow: image recognition and dataset categorization*, pp. 63–72, Springer, 2021.
- [42] G. Huang, Z. Liu, L. Van Der Maaten, and K. Weinberger, "Densely connected convolutional networks," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, pp. 4700–4708, 2017.
- [43] F. Chollet, "Xception: Deep learning with depthwise separable convolutions," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, pp. 1251–1258, 2017.
- [44] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, and S. Gelly, *et al.*, "An image is worth 16x16 words: Transformers for image recognition at scale," *arXiv preprint arXiv:2010.11929*, 2020.
- [45] K. Vinogradova, A. Dibrov, and G. Myers, "Towards interpretable semantic segmentation via gradient-weighted class activation mapping (student abstract)," in *Proc. AAAI Conf. Artif. Intell.*, 2020.

- 2088 [46] Y. Zhang, Y. Zhu, J. Liu, and W. Yu, “An interpretability
2089 optimization method for deep learning networks based on
2090 grad-cam,” *IEEE*, 2024.
- 2091 [47] N. Makke and S. Chawla, “Interpretable scientific discov-
2092 ery with symbolic regression: a review,” *Applied Intelli-*
2093 *gence Review*, 2024.
- 2094 [48] B. Ghojogh, A. Ghodsi, F. Karray, and M. Crowley,
2095 “Uniform manifold approximation and projection (umap)
2096 and its variants: tutorial and survey,” *arXiv preprint*
2097 *arXiv:2109.02508*, 2021.
- 2098 [49] S. Li, X. Wang, and T. Rasing, “Magnetic skyrmions: Basic
2099 properties and potential applications,” *Interdiscip. Mater.*,
2100 vol. 2, no. 2, pp. 260–289, 2023.
- 2101 [50] C. Fillion, J. Fischer, and R. Kumar, *et al.*, “Gate-controlled
2102 skyrmion and domain wall chirality,” *Nat. Commun.*,
2103 vol. 13, no. 1, p. 5257, 2022.
- 2104 [51] I. P. Miranda, A. Klautau, A. Bergman, and H. Petrilli,
2105 “Band filling effects on the emergence of magnetic
2106 skyrmions: Pd/Fe and Pd/Co bilayers on Ir(111),” *Phys.*
2107 *Rev. B*, vol. 105, p. 224413, 2022.
- 2108 [52] F. Donati and A. Heinrich, “A perspective on surface-
2109 adsorbed single atom magnets as atomic-scale magnetic
2110 memory,” *Appl. Phys. Lett.*, vol. 119, no. 16, p. 160503,
2111 2021.
- 2112 [53] M. T. Ribeiro, S. Singh, and C. Guestrin, ““Why should I
2113 trust you?” Explaining the predictions of any classifier,” in
2114 *Proc. 22nd ACM SIGKDD Int. Conf. Knowl. Discov. Data*
2115 *Min. (KDD)*, pp. 1135–1144, 2016.
- 2116 [54] S. M. Lundberg and S.-I. Lee, “A unified approach to
2117 interpreting model predictions,” in *Adv. Neural Inf. Process.*
2118 *Syst. (NeurIPS)*, vol. 30, 2017.
- 2119 [55] A. Gorji, A. Amrollahi, and A. Krause, “SHAP values via
2120 sparse Fourier representation,” in *Adv. Neural Inf. Process.*
2121 *Syst. (NeurIPS)*, vol. 38, 2025.
- 2122 [56] W. Samek, A. Binder, G. Montavon, S. Lapuschkin, and
2123 K.-R. Müller, “Evaluating the visualization of what a deep
2124 neural network has learned,” *IEEE Trans. Neural Netw.*
2125 *Learn. Syst.*, vol. 28, no. 11, pp. 2660–2673, 2017.